

HDF5 Dynamically Loaded Filters

The HDF Group

This document describes the HDF5 dynamically loaded filter feature (HDF5 filter plugin) that will be released in the HDF5 version 1.8.11 in May 2013.

This feature allows an HDF5 application to apply a non-native HDF5 filter during I/O operations at runtime. The application does not need to register the filter with the HDF5 Library. There is no need for adding the filter to the HDF5 Library and then linking the application with the modified library.

For example, the HDF5 1.8.11 command-line tools such as h5dump or h5ls downloaded from The HDF Group website, or Java-based browsing tool HDFView that uses 1.8.11 will be able to read LZF or bzip2 compressed datasets created by H5Py or PyTables software.

The dynamically loaded filter feature was sponsored by **Deutsches Elektronen-Synchrotron (DESY)**, Hamburg, Germany.

Contents

1	Introduction	3
2	Requirements	5
2.1	Functional Requirements	5
2.2	Non-functional Requirements	5
3	Programming Model for Applications	6
3.1	Applying a Third-party Filter When Creating and Writing a Dataset.....	6
3.2	Reading Data with an Applied Third-party Filter	6
3.3	A Word of Caution When Using Custom Filters	8
4	Programming Model for HDF5 Filter Plugins.....	10
4.1	Writing a Filter Function.....	10
4.2	Registering a Filter with The HDF Group	10
4.3	Creating an HDF5 Filter Plugin.....	11
4.4	Installing an HDF5 Filter Plugin.....	12
5	Design.....	13
5.1	Data Writing	13
5.2	Data Reading	14
5.3	New HDF5 Source Files for Dynamically Loaded Filters	15
5.4	Testing Dynamically Loaded Filters	15
6	Proposed Changes to the HDF5 Command-line Tools.....	16
6.1	h5dump, h5ls, and h5copy	16
6.2	h5repack.....	16
7	Building an HDF5 bzip2 Plugin Example	17
8	Example: Writing and Reading Compressed Data.....	18
9	HDF5 bzip2 Filter Plugin	22
10	Example of Makefile.....	26
11	Revision History.....	27
12	References.....	28

1 Introduction

The HDF5 Library provides an internal mechanism to compress an HDF5 dataset's raw data and the links stored in an HDF5 group object. Readers who are unfamiliar with HDF5 compression are encouraged to view the HDF5 tutorials, the *HDF5 User's Guide*, the *HDF5 Reference Manual* and other documentation available from the <http://www.hdfgroup.org> website.

The HDF5 compression mechanism is a part of the HDF5 “filter pipeline” that modifies an application's data during the I/O operations. The pipeline was designed to be extensible. New filters could be easily added to the pipeline by an application, or the HDF5 Library code can be modified with a new filter to be later used by an application.

The HDF5 Library supports several “internal” filters (see in the *HDF5 User's Guide*, section 5.4.2, “Data Pipeline Filters”). Two filters, *deflate* and *gzip*, depend on third-party compression libraries. All internal filters are configurable. They can be added or removed during the configuration step when the HDF5 Library is built.

While currently available HDF5 “internal” compression methods work reasonably well on users' datasets, there are certain drawbacks to the current implementation. First, the “internal” compression methods may not provide the optimal compression ratio, as do some newly developed or specialized compression methods. Secondly, if a data provider wants to use a “non-internal” compression for storing the data in HDF5, he/she has to write a filter function that uses the new compression method and then register it with the library. Data consumers of such HDF5 files will need to have the new filter function and use it with their applications to read the data, or they will need a modified version of the HDF5 Library that has the new filter as a part of the library.

If a user of such data does not have a modified HDF5 Library installed on his system, command-line tools such as *h5dump* or *h5ls* will not be able to display the compressed data. Furthermore, it would be practically impossible to determine the compression method used, making the data stored in HDF5 useless.

It is clear that the current HDF5 filter mechanism, while extensible, does not work well with third-party filters. It would be a maintenance nightmare to keep adding and supporting new compression methods in HDF5. For any set of HDF5 “internal” filters, there always will be data with which the “internal” filters will not achieve the optimal performance needed to address data I/O and storage problems. Thus the current HDF5 filter mechanism should be enhanced to address the issues discussed above.

We have added a new feature to HDF5 called “dynamically loaded filters in HDF5.” This feature will make the HDF5 third-party filters available to an application at runtime. The third-party HDF5 filter function has to be a part of the HDF5 filter plugin installed on the system as a shared library or DLL.

To use a third-party filter an HDF5 application should call the `H5Pset_filter` function when setting the filter pipeline for a dataset creation property. The HDF5 Library will register the filter with the library and the filter will be applied when data is written to the file.

When an application reads data compressed with a third-party HDF5 filter, the HDF5 Library will search for the required filter plugin, register the filter with the library (if the filter function is not registered) and apply it to the data on the read operation.

This document describes this new feature and provides details in the following sections.

2 Requirements

The current design for the HDF5 dynamically loaded filters feature addresses the requirements listed in the sections below.

2.1 Functional Requirements

1. There are **no changes to the HDF5 file format**.
2. There are **no changes to the HDF5 interface and the programming model** (see the “Programming Model for Applications” section on page 6) in order to use the feature.
3. The HDF5 third-party filters are available as **shared libraries or DLLs** on the user’s system. The minimum content of the shared library satisfies the requirements as documented in the “Programming Model for HDF5 Filter Plugins” section on page 10.
4. There are predefined **default locations** where the HDF5 Library searches the shared libraries or DLLs with the HDF5 filter functions.
5. The **default location may be overwritten** by an environment variable.
6. Once a filter plugin library is loaded, it stays loaded until the HDF5 Library is closed.

Future enhancements:

- The default search locations can be specified in a configuration file.

2.2 Non-functional Requirements

1. The HDF Group maintains a list of the third-party filters registered with HDF5 to facilitate tracking of a filter provenance. The current procedure is outlined at <http://www.hdfgroup.org/services/contributions.html>.
2. The HDF group publishes a specification and provides examples of dynamic loaded filters libraries.

Future enhancements:

- The HDF Group provides a repository such as SVN or GIT for the third-party filters and outlines a testing procedure that can be used by the filters maintainers to test the filters with the HDF5 libraries under development.

3 Programming Model for Applications

This section describes the programming model for an application that uses a third-party HDF5 filter plugin to write or read data. For simplicity of presentation, it is assumed that the HDF5 filter plugin is available on the system in a default location. The HDF5 filter plugin is discussed in detail in the “Programming Model for HDF5 Filter Plugins” section on page 10.

3.1 Applying a Third-party Filter When Creating and Writing a Dataset

A third-party filter can be added to the HDF5 filter pipeline by using the `H5Pset_filter` function, as a user would do in the past. The identification number and the filter parameters should be available to the application. For example, if the application intends to apply the HDF5 bzip2 compression filter that was registered with The HDF Group and has an identification number 307 (see <http://www.hdfgroup.org/services/contributions.html>) then the application would follow the steps as outlined below:

```
dcpl = H5Pcreate (H5P_DATASET_CREATE);
status = H5Pset_filter (dcpl, (H5Z_filter_t)307, H5Z_FLAG_MANDATORY,
                        (size_t)6, cd_values);

dset = H5Dcreate (file, DATASET, H5T_STD_I32LE, space, H5P_DEFAULT, dcpl,
status = H5Dwrite (dset, H5T_NATIVE_INT, H5S_ALL, H5S_ALL, H5P_DEFAULT,
                    wdata[0]);
```

3.2 Reading Data with an Applied Third-party Filter

An application does not need to do anything special to read the data with a third-party filter applied. For example, if one wants to read data written in 3.1 the following regular steps should be taken:

```
file = H5Fopen (FILE, H5F_ACC_RDONLY, H5P_DEFAULT);
dset = H5Dopen (file, DATASET, H5P_DEFAULT);
H5Dread (dset, H5T_NATIVE_INT, H5S_ALL, H5S_ALL, H5P_DEFAULT,
        rdata[0]);
```

The 1.8.11 version of the command-line utility `h5dump`, for example, will read and display the data as shown:

```
[epourmal@jam example]$ ./h5dump -p *.h5
HDF5 "h5ex_d_bzip2.h5" {
GROUP "/" {
  DATASET "DS1" {
    DATATYPE  H5T_STD_I32LE
    DATASPACE  SIMPLE { ( 32, 64 ) / ( 32, 64 ) }
    STORAGE_LAYOUT {
      CHUNKED ( 4, 8 )
      SIZE 6410 (1.278:1 COMPRESSION)
    }
    FILTERS {
```

```

    USER_DEFINED_FILTER {
        FILTER_ID 307
        COMMENT bzip2
        PARAMS { 6 }
    }
}
...
}
DATA {
(0,0): 0, -1, -2, -3, -4, -5, -6, -7, -8, -9, -10, -11, -12, -13, -14,
(0,15): -15, -16, -17, -18, -19, -20, -21, -22, -23, -24, -25, -26,
(0,27): -27, -28, -29, -30, -31, -32, -33, -34, -35, -36, -37, -38,
(0,39): -39, -40, -41, -42, -43, -44, -45, -46, -47, -48, -49, -50,
(0,51): -51, -52, -53, -54, -55, -56, -57, -58, -59, -60, -61, -62,
(0,63): -63,
.....

```

Compare the output above with the output of the 1.8.10 version of h5dump shown below:

```

[epourmal@jam example]$ h5dump -p *.h5
HDF5 "h5ex_d_bzip2.h5" {
GROUP "/" {
  DATASET "DS1" {
    DATATYPE  H5T_STD_I32LE
    DATASPACE  SIMPLE { ( 32, 64 ) / ( 32, 64 ) }
    STORAGE_LAYOUT {
      CHUNKED ( 4, 8 )
      SIZE 6410
    }
    FILTERS {
      UNKNOWN_FILTER {
        FILTER_ID 307
        COMMENT HDF5 bzip2 filter; see http://www.hdfgroup.org/services/contributions.html
        PARAMS { 6 }
      }
    }
  }
  ...
}
  DATA {h5dump error: unable to print data
}
}
}
}

```

The h5dump utility version 1.8.10 cannot read data compressed with bzip2 because the filter is not available to the HDF5 library.

The complete example of writing and reading data with an applied third-party HDF5 filter is shown in the “Example: Writing and Reading Compressed Data” section on page 18. Please notice that the application DOES NOT need to link with the HDF5 plugin library.

3.3 A Word of Caution When Using Custom Filters

Data goes through the HDF5 filter pipeline only when it is written to the file or read into application memory space from the file. For example, the I/O operation is triggered with a call to `H5Fflush`, or when a data item (HDF5 metadata or a raw data chunk) is evicted from the cache or brought into the cache. Please notice that `H5Dread/write` calls on the chunked datasets do not necessarily trigger I/O since the HDF5 Library uses a separate chunk cache.

A data item may remain in the cache until the HDF5 Library is closed. If the HDF5 plugin that has to be applied to the data item becomes unavailable before the file and all objects in the file are closed, an error will occur. The following example demonstrates the issue. Please notice the position of the `H5Zunregister` call:

```
/*
 * Create a new group using compression.
 */
gcpl = H5Pcreate (H5P_GROUP_CREATE);
status = H5Pset_filter(gcpl,H5Z_FILTER_BZIP2,H5Z_FLAG_MANDATORY,(size_t)1, cd_values);
group = H5Gcreate (file, GNAME, H5P_DEFAULT, gcpl, H5P_DEFAULT);
for (i=0; i < NGROUPS; i++) {
    sprintf(name, "group_%d", i);
    tmp_id = H5Gcreate (group, name, H5P_DEFAULT, H5P_DEFAULT, H5P_DEFAULT);
    status = H5Gclose(tmp_id);
}

status = H5Pclose (gcpl);
status = H5Gclose (group);
/*
 * Unregister the filter. Call to H5Fclose will fail because the library tries
 * to apply the filter that is not available anymore. This has a cascade effect
 * on H5Fclose.
 */
H5Zunregister(H5Z_FILTER_BZIP2);
status = H5Fclose (file);
```

Here is an error stack produced by the program:

```
HDF5-DIAG: Error detected in HDF5 (1.9.149) thread 0:
#000: H5F.c line 2060 in H5Fclose(): decrementing file ID failed
  major: Object atom
  minor: Unable to close file
#001: H5I.c line 1406 in H5I_dec_app_ref(): can't decrement ID ref count
  major: Object atom
  minor: Unable to decrement reference count
#002: H5F.c line 1837 in H5F_close(): can't close file
  major: File accessibility
  minor: Unable to close file
.....
#026: H5Z.c line 1045 in H5Z_find(): required filter is not registered
  major: Data filters
  minor: Object not found
```


To avoid the problem make sure to close all objects to which the filter is applied and flush them using the `H5Fflush` call before unregistering the filter.

4 Programming Model for HDF5 Filter Plugins

This section describes how to create an HDF5 filter, an HDF5 filter plugin, and how to install the HDF5 plugin on the system.

4.1 Writing a Filter Function

The HDF5 filter function for the dynamically loaded filter feature should be written as any custom filter described in <http://www.hdfgroup.org/HDF5/doc/H5.user/Filters.html>. See the “Example” section, section 5, of that document to get an idea of the simple filter function, and see the example of the more sophisticated HDF5 bzip2 filter function in the “Building an HDF5 bzip2 Plugin Example” section on page 17. The HDF5 bzip2 filter function is also available for download from https://svn.hdfgroup.uiuc.edu/hdf5_plugins/trunk/BZIP2.

The user has to remember a few things when writing an HDF5 filter function.

1. An HDF5 filter is bidirectional.

The filter handles both input and output to the file; a flag is passed to the filter to indicate the direction.

2. An HDF5 filter operates on a buffer.

The filter reads data from a buffer, performs some sort of transformation on the data, places the result in the same or new buffer, and returns the buffer pointer and size to the caller.

3. An HDF5 filter should return zero in the case of failure.

The signature of the HDF5 filter function and the accompanying filter structure (see the section below) are described in the *HDF5 Reference Manual*

http://www.hdfgroup.org/HDF5/doc/RM/RM_H5Z.html#Compression-Register.

4.2 Registering a Filter with The HDF Group

If you are writing a filter that will be used by others, it would be a good idea to request a filter identification number and register it with The HDF Group. Please follow the procedure described at <http://www.hdfgroup.org/services/contributions.html>.

The HDF Group anticipates that developers of HDF5 filter plugins will not only register new filters, but will also provide links to the source code and/or binaries for the corresponding HDF5 filter plugins.

It is very important for the users of the filter that developers provide filter information in the “name” field of the filter structure, for example:

```
const H5Z_class2_t H5Z_BZIP2[1] = {
    H5Z_CLASS_T_VERS,      /* H5Z_class_t version */
    (H5Z_filter_t) 307,    /* Filter id number */
    1,                     /* encoder_present flag (set to true) */
    1,                     /* decoder_present flag (set to true) */
    "HDF5 bzip2 filter; see
```

```

    http://www.hdfgroup.org/services/contributions.html",
                                /* Filter info */
    NULL,                        /* The "can apply" callback */
    NULL,                        /* The "set local" callback */
    (H5Z_func_t) H5Z_filter_bzip2, /* The filter function */
};

```

The HDF5 Library and command-line tools have access to the “name” field. An application can use the `H5Pget_filter<*>` functions to retrieve information about the filters.

Using the example of the structure above, the `h5dump` tool will print the string “HDF5 bzip2 filter found at ...” pointing users to the applied filter (see the example in the “Reading Data with an Applied Third-party Filter” section on page 6) thus solving the problem of the filter’s origin.

4.3 Creating an HDF5 Filter Plugin

The HDF5 filter plugin source should include:

1. The `H5PLextern.h` header file from the HDF5 distribution.
2. The definition of the filter structure (see the example shown in the section above).
3. The filter function (for example, `H5Z_filter_bzip2`).
4. The two functions necessary for the HDF5 Library to find the correct type of the plugin library while loading it at runtime and to get information about the filter function:

```

H5PL_type_t  H5PLget_plugin_type(void);
const void*  H5PLget_plugin_info(void);

```

Here is an example of the functions above for the HDF5 bzip2 filter:

```

H5PL_type_t  H5PLget_plugin_type(void) {return H5PL_TYPE_FILTER;}
const void*  H5PLget_plugin_info(void) {return H5Z_BZIP2;}

```

5. Other functions such as the source of the compression library may also be included.

Build the HDF5 filter plugin as a shared library. The following steps should be taken:

1. When compiling, point to the HDF5 header files.
2. Use the appropriate linking flags.
3. Link with any required external libraries.
4. For example, if `libbz2.so` is installed on a Linux system, the HDF5 bzip2 plugin library `libH5Zbzip2.so` may be linked with `libbz2.so` instead of including bzip2 source into the plugin library.

The complete example of the HDF5 bzip2 plugin library is provided at http://svn.hdfgroup.uiuc.edu/hdf5_plugins/trunk/BZIP2/ and can be adopted for other

plugins. A simple Makefile file for Linux systems is shown in the “Example of Makefile” section on page 26.

4.4 Installing an HDF5 Filter Plugin

The default directory for an HDF5 filter plugin library is defined on UNIX-like systems as

```
"/usr/local/hdf5/lib/plugin"
```

and on Windows systems as

```
"%ALLUSERSPROFILE%/hdf5/lib/plugin".
```

The default path can be overwritten by a user with the `HDF5_PLUGIN_PATH` environment variable. Several directories can be specified for the search path using “:” as a path separator for UNIX-like systems and “;” for Windows.

Readers are encouraged to try the example in the “Building an HDF5 bzip2 Plugin Example” section on page 17.

5 Design

Dynamic loading of the HDF5 filter plugin (or filter library) is triggered only by two events: when an application calls the `H5Pset_filter` function to set the filter for the first time, or when the data to which the filter is applied is read for the first time.

5.1 Data Writing

This section outlines the usage of the new feature on the write operation.

When the application enables a filter with `H5Pset_filter`, the library will search for it by its filter identifier. The library will first search for the filter among its internal filters and custom filters registered with the library, then it will search in the list of filters that have been dynamically loaded. If it cannot find the filter, it will search for the filter library in the default path or the path indicated by the environment variable `HDF5_PLUGIN_PATH`.

When found, the HDF5 Library will load the filter library, register the filter, and add it to the filter pipeline. There is no need for the application to register the dynamic filter with `H5Zregister`, as this step is done internally in `H5Pset_filter`. The `H5PL_load` function in the `H5Pocpl.c` file in the `src` directory of the HDF5 Library distribution does the searching and loading of the filter library.

Once the filter library is loaded, it will stay loaded until the HDF5 Library is closed or until the filter is unregistered with the `H5Zunregister` call. The `H5Zregister` function is still required when a third-party filter is linked into an application and is not provided as a shared library.

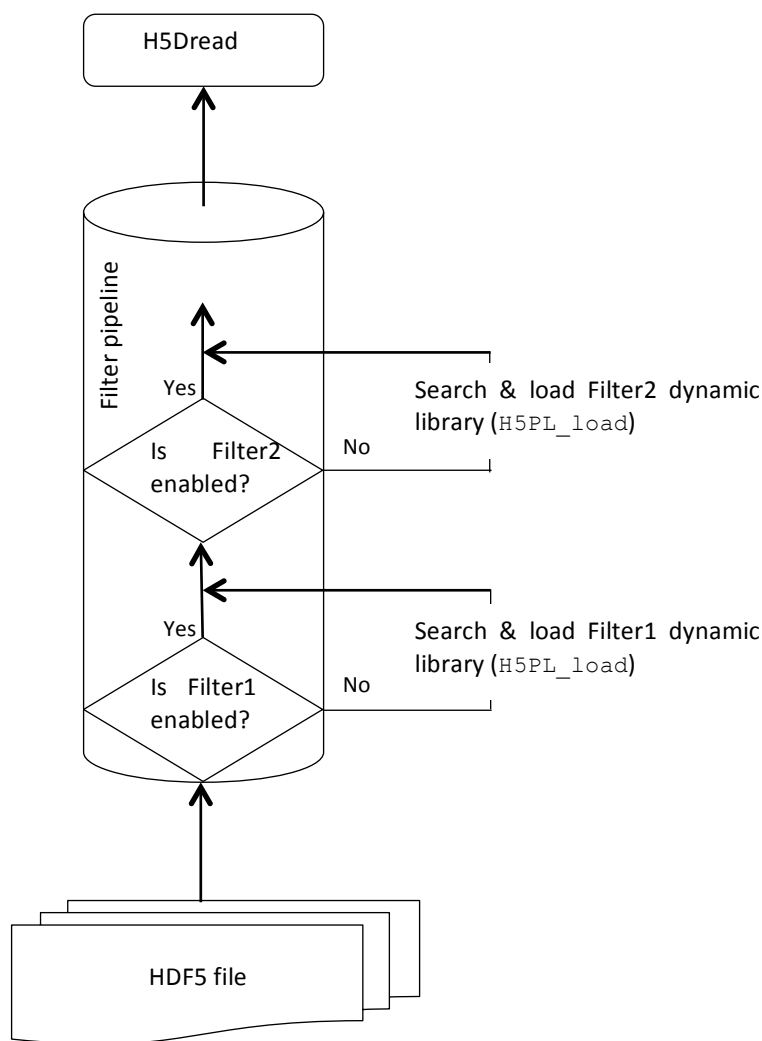
If the dynamically loaded filter has any auxiliary data, such as the compression level for `bzip2` filter, it can be passed in through the `cd_values` parameter of `H5Pset_filter`.

5.2 Data Reading

This section outlines the usage of the new feature on the read operation.

At the application level, the user does not need to do anything special. Inside the HDF5 Library, when the library sees a filter identification value associated with a data item, it first searches the filter among its internal filters, and then it searches in the list of dynamically loaded filters that have been loaded in the library. If it cannot find the filter, it will search for and load the filter plugin library. Once the filter library is loaded, it will stay open until the HDF5 Library is closed. Below is a simple diagram showing the design of the library.

Figure 1. Reading data with a dynamic loaded filter



The required changes were implemented in the `H5Z_pipeline` function that resides in the `H5Z.c` file in the `src` directory of the HDF5 Library distribution.

5.3 New HDF5 Source Files for Dynamically Loaded Filters

Three new source files were added to the HDF5 source code distribution under the `src` directory: `H5PLextern.h`, `H5PLprivate.h`, and `H5PL.c`. The `H5PL.c` file contains functions for managing plugins.

5.4 Testing Dynamically Loaded Filters

The `plugin.c` file in the `test` directory contains a test for the feature. The test is controlled by the `test_plugin.sh.in` script in the same directory.

The HDF Group will use `bzip2` (and other plugins if they become available) to test the feature with the HDF5 libraries under development. For more on the `bzip2` plugin, see the “Building an HDF5 `bzip2` Plugin Example” section on page 17.

6 Proposed Changes to the HDF5 Command-line Tools

6.1 h5dump, h5ls, and h5copy

Since reading the data is transparent to applications including the HDF5 command-line tools, these command line tools do not require any changes. A small improvement was made to h5dump output to replace “UNKNOWN_FILTER” with “USER_DEFINED_FILTER”.

6.2 h5repack

As in the case of the tools mentioned in the section above, h5repack does not require any changes for reading a dataset with an applied third-party filter.

To support an application’s use of a third-party filter while repacking datasets, new syntax is needed to provide filter information to the tool. The required information is the filter identification number, the number of the filter parameters, and the values of the filter parameters. The proposed syntax to specify a user defined filter and its parameters is

```
h5repack -f UD={ID:k; N:m; CD_VAL:[n1,...,nm]}... .
```

where *k* is a filter identifier, *m* is a number of values in the CD_VAL array, and *n₁,...,n_m* are compression parameters. For example, to use bzip2 compression with h5repack, one would use

```
h5repack -f UD={ID:307; N:1; CD_VAL:[9]} file1.h5 file2.h5
```

The enhancements to h5repack will be implemented in the HDF5 release 1.8.12.

7 Building an HDF5 bzip2 Plugin Example

The HDF Group provides an example of the HDF5 filter plugin that can be checked out from https://svn.hdfgroup.uiuc.edu/hdf5_plugins/trunk/BZIP2. It contains the source code for the bzip2 plugin library and an example that uses the plugin. It requires the HDF5 Library with the dynamically loaded feature and the bzip2 library being available on the system.

The plugin and the example can be built using configure or CMake commands. For instructions on how to build with CMake, see the README.txt file in the source code distribution. The bzip2 library that can be built with CMake is available from <https://svn.hdfgroup.uiuc.edu/bzip2/>.

Please follow the instructions below to build with configure.

1. Obtain, build, and install the HDF5 Library with the dynamically loaded filter feature.

You may checkout the source from http://svn.hdfgroup.uiuc.edu/hdf5/branches/hdf5_1_8/ or download <http://www.hdfgroup.uiuc.edu/ftp/pub/outgoing/hdf5/snapshots/v18/hdf5-1.8.11-snap16.tar.gz>.

2. Make sure that libbz2.so* is installed on your system. On most UNIX-like systems it will be under /usr/lib.
3. Obtain the source code for the plugin using SVN, or download via FTP the source tar ball:

```
svn export https://svn.hdfgroup.uiuc.edu/hdf5\_plugins/trunk/BZIP2 <plugin-dir>  
ftp://ftp.hdfgroup.uiuc.edu/pub/outgoing/HDF5-plugins/BZIP2-plugin.tar.gz
```

4. You can build in place, or you can use the “srdir” configure option. If you build in place as shown below, you will find the plugin under the <plugin-dir>/plugins directory.

```
cd <plugin-dir>  
./configure --with-hdf5=<INSTALL-HDF5-DIR> --with-bz2lib=/usr  
make  
make check  
make install
```

5. Setup the environment variable HDF5_PLUGIN_PATH to point to the <plugin-dir>/plugins directory with the command line shown below:

```
setenv HDF5_PLUGIN_PATH <plugin-dir>/plugins
```

6. Use <INSTALL-HDF5-DIR>/bin/h5dump with the h5ex_d_bzip2.h5 file found under the example directory of the bzip2 plugin distribution to get an output as shown in the “Reading Data with an Applied Third-party Filter” section on page 6.

8 Example: Writing and Reading Compressed Data

This example shows how to write and read data compressed with the bzip2 compression.

```

1/*****
2
3 This example shows how to write data and read it from a dataset
4 using bzip2 compression.
5 bzip2 filter is not available in HDF5.
6 The example uses a new feature available in HDF5 version 1.8.11
7 to discover, load and register filters at runtime.
8
9 *****/
10
11#include "hdf5.h"
12#include <stdio.h>
13#include <stdlib.h>
14
15#define FILE          "h5ex_d_bzip2.h5"
16#define DATASET       "DS1"
17#define DIM0          32
18#define DIM1          64
19#define CHUNK0        4
20#define CHUNK1        8
21
22int
23main (void)
24{
25     hid_t          file, space, dset, dcpl;    /* Handles */
26     herr_t          status;
27     H5Z_filter_t     filter_id = 0;
28     char            filter_name[80];
29     hsize_t          dims[2] = {DIM0, DIM1},
30                     chunk[2] = {CHUNK0, CHUNK1};
31     size_t           nelmts = 1;               /* number of elements in cd_values */
32     const unsigned   int cd_values[1] = {6};   /* bzip2 default level is 2 */
33     unsigned int     values_out[1] = {99};
34     unsigned int     flags;
35     htri_t           avail;
36     unsigned         filter_config;
37     int              wdata[DIM0][DIM1],       /* Write buffer */
38                     rdata[DIM0][DIM1],       /* Read buffer */
39                     max,
40                     i, j;
41
42     /*
43      * Initialize data.
44      */
45     for (i=0; i<DIM0; i++)
46         for (j=0; j<DIM1; j++)
47             wdata[i][j] = i * j - j;
48

```

```
49  /*
50   * Create a new file using the default properties.
51   */
52  file = H5Fcreate (FILE, H5F_ACC_TRUNC, H5P_DEFAULT, H5P_DEFAULT);
53
54  /*
55   * Create dataspace. Setting maximum size to NULL sets the maximum
56   * size to be the current size.
57   */
58  space = H5Screate_simple (2, dims, NULL);
59
60  /*
61   * Create the dataset creation property list, add the bzip2
62   * compression filter and set the chunk size.
63   */
64  dcpl = H5Pcreate (H5P_DATASET_CREATE);
65  status = H5Pset_filter (dcpl, H5Z_FILTER_BZIP2, H5Z_FLAG_MANDATORY,
66                          (size_t)6, cd_values);
67
68  /*
69   * Check that filter is registered with the library now.
70   * If it is registered, retrieve filter's configuration.
71   */
72  avail = H5Zfilter_avail(H5Z_FILTER_BZIP2);
73  if (avail) {
74      status = H5Zget_filter_info (H5Z_FILTER_BZIP2, &filter_config);
75      if ( (filter_config & H5Z_FILTER_CONFIG_ENCODE_ENABLED) &&
76          (filter_config & H5Z_FILTER_CONFIG_DECODE_ENABLED) )
77          printf ("bzip2 filter is available for encoding and decoding.\n");
78  }
79  status = H5Pset_chunk (dcpl, 2, chunk);
80
81  /*
82   * Create the dataset.
83   */
84  printf ("....Writing bzip2 compressed data .....\\n");
85  dset = H5Dcreate (file, DATASET, H5T_STD_I32LE, space, H5P_DEFAULT, dcpl,
86                  H5P_DEFAULT);
87
88  /*
89   * Write the data to the dataset.
90   */
91  status = H5Dwrite (dset, H5T_NATIVE_INT, H5S_ALL, H5S_ALL, H5P_DEFAULT,
92                    wdata[0]);
93
94  /*
95   * Close and release resources.
96   */
97  status = H5Pclose (dcpl);
98  status = H5Dclose (dset);
99  status = H5Sclose (space);
100 status = H5Fclose (file);
101 status = H5close();
```

```

101
102
103     printf ("....Close the file and reopen for reading .....\\n");
104     /*
105      * Now we begin the read section of this example.
106      */
107
108     /*
109      * Open file and dataset using the default properties.
110      */
111     file = H5Fopen (FILE, H5F_ACC_RDONLY, H5P_DEFAULT);
112     dset = H5Dopen (file, DATASET, H5P_DEFAULT);
113
114     /*
115      * Retrieve dataset creation property list.
116      */
117     dcpl = H5Dget_create_plist (dset);
118
119     /*
120      * Retrieve and print the filter id, compression level and filter's name
for bzip2.
121      */
122     filter_id = H5Pget_filter2 (dcpl, (unsigned) 0, &flags, &nelmts,
                                values_out, sizeof(filter_name), filter_name, NULL);
123     printf ("Filter info is available from the dataset creation property \\n ");
124     printf ("  Filter identifier is ");
125     switch (filter_id) {
126         case H5Z_FILTER_BZIP2:
127             printf ("%d\\n", filter_id);
128             printf ("    Number of parameters is %d with the value %u\\n",
                                nelmts, values_out[0]);
129             printf ("    To find more about the filter check %s\\n",
                                filter_name);
130             break;
131         default:
132             printf ("Not expected filter\\n");
133             break;
134     }
135     /*
136      * Check that filter is not registered with the library yet.
137      */
138     avail = H5Zfilter_avail(H5Z_FILTER_BZIP2);
139     if (!avail)
140         printf ("bzip2 filter is not yet available after the H5Pget_filter
                                call.\\n");
141     else
142         return 1;
143
144
145     /*
146      * Read the data using the default properties.
147      */
148     printf ("....Reading bzip2 compressed data .....\\n");

```

```
149     status = H5Dread (dset, H5T_NATIVE_INT, H5S_ALL, H5S_ALL, H5P_DEFAULT,
150                       rdata[0]);
151
152     /*
153     * Find the maximum value in the dataset, to verify that it was
154     * read correctly.
155     */
156     max = rdata[0][0];
157     for (i=0; i<DIM0; i++)
158         for (j=0; j<DIM1; j++) {
159             /*printf("%d \n", rdata[i][j]); */
160             if (max < rdata[i][j])
161                 max = rdata[i][j];
162         }
163     /*
164     * Print the maximum value.
165     */
166     printf ("Maximum value in %s is %d\n", DATASET, max);
167     /*
168     * Check that filter is registered with the library now.
169     */
170     avail = H5Zfilter_avail(H5Z_FILTER_BZIP2);
171     if (avail)
172         printf ("bzip2 filter is available now since H5Dread triggered loading
173                 of the filter.\n");
174
175     /*
176     * Close and release resources.
177     */
178     status = H5Pclose (dcpl);
179     status = H5Dclose (dset);
180     status = H5Fclose (file);
181
182     return 0;
183 }
```

9 HDF5 bzip2 Filter Plugin

This section provides partial code for the bzip2 filter plugin. For complete code, check http://svn.hdfgroup.uiuc.edu/hdf5_plugins/trunk/BZIP2/src/H5Zbzip2.c

```

1  /*
2  * This file is an example of an HDF5 filter plugin.
3  * The filter function H5Z_filter_bzip2 was adopted from
4  * PyTables http://www.pytables.org.
5  * The plugin can be used with the HDF5 Library version 1.8.11 to read
6  * HDF5 datasets compressed with bzip2 created by PyTables.
7  */
8
9  /*
10 *
11 Copyright Notice and Statement for PyTables Software Library and Utilities:
12
...
45 */
46 #include <sys/types.h>
47 #include <stdlib.h>
48 #include <string.h>
49 #include <assert.h>
50 #include <stdio.h>
51 #include <H5PExtern.h>
52
53 #include <bzlib.h>
54
55 static size_t H5Z_filter_bzip2(unsigned int flags, size_t cd_nelmts,
56                               const unsigned int cd_values[], size_t nbytes,
57                               size_t *buf_size, void **buf);
58
59 /*
...
62 * See http://www.hdfgroup.org/services/contributions.html for more information.
63 *
64 * If you intend your plugin to be used by others, please register your filter
65 * with The HDF Group.
66 */
67 #define H5Z_FILTER_BZIP2 307
68
69 const H5Z_class2_t H5Z_BZIP2[1] = {{
70     H5Z_CLASS_T_VERS, /* H5Z_class_t version */
71     (H5Z_filter_t)H5Z_FILTER_BZIP2, /* Filter id number */
72     1, /* encoder_present flag (set to true) */
73     1, /* decoder_present flag (set to true) */
74     "HDF5 bzip2 filter; see http://www.hdfgroup.org/services/contributions.html",
75     /* Filter name for debugging */
76     NULL, /* The "can apply" callback */
77     NULL, /* The "set local" callback */
78     (H5Z_func_t)H5Z_filter_bzip2, /* The actual filter function */
79 }};
80

```

```

81 H5PL_type_t    H5PLget_plugin_type(void) {return H5PL_TYPE_FILTER;}
82 const void *H5PLget_plugin_info(void) {return H5Z_BZIP2;}
83
84 static size_t H5Z_filter_bzip2(unsigned int flags, size_t cd_nelmts,
85                                const unsigned int cd_values[], size_t nbytes,
86                                size_t *buf_size, void **buf)
87 {
88     char *outbuf = NULL;
89     size_t outbuflen, outdatalen;
90     int ret;
91
92     if (flags & H5Z_FLAG_REVERSE) {
93
94         /** Decompress data.
95          **
96
97         ...
98
99         **/
100
101         bz_stream stream;
102         char *newbuf = NULL;
103         size_t newbuflen;
104
105         /* Prepare the output buffer. */
106         outbuflen = nbytes * 3 + 1; /* average bzip2 compression ratio is 3:1 */
107         outbuf = malloc(outbuflen);
108         if (outbuf == NULL) {
109             fprintf(stderr, "memory allocation failed for bzip2 decompression\n");
110             goto cleanupAndFail;
111         }
112
113         /* Use standard malloc()/free() for internal memory handling. */
114         stream.bzalloc = NULL;
115         stream.bzfree = NULL;
116         stream.opaque = NULL;
117
118         /* Start decompression. */
119         ret = BZ2_bzDecompressInit(&stream, 0, 0);
120         if (ret != BZ_OK) {
121             fprintf(stderr, "bzip2 decompression start failed with error %d\n", ret);
122             goto cleanupAndFail;
123         }
124
125         /* Feed data to the decompression process and get decompressed data. */
126         stream.next_out = outbuf;
127         stream.avail_out = outbuflen;
128         stream.next_in = *buf;
129         stream.avail_in = nbytes;
130         do {
131             ret = BZ2_bzDecompress(&stream);
132             if (ret < 0) {
133                 fprintf(stderr, "BUG: bzip2 decompression failed with error %d\n", ret);
134                 goto cleanupAndFail;
135             }
136
137             if (ret != BZ_STREAM_END && stream.avail_out == 0) {
138                 /* Grow the output buffer. */

```

```

142     newbuflen = outbuflen * 2;
143     newbuf = realloc(outbuf, newbuflen);
144     if (newbuf == NULL) {
145         fprintf(stderr, "memory allocation failed for bzip2 decompression\n");
146         goto cleanupAndFail;
147     }
148     stream.next_out = newbuf + outbuflen; /* half the new buffer behind */
149     stream.avail_out = outbuflen; /* half the new buffer ahead */
150     outbuf = newbuf;
151     outbuflen = newbuflen;
152 }
153 } while (ret != BZ_STREAM_END);
154
155 /* End compression. */
156 outdatalen = stream.total_out_lo32;
157 ret = BZ2_bzDecompressEnd(&stream);
158 if (ret != BZ_OK) {
159     fprintf(stderr, "bzip2 compression end failed with error %d\n", ret);
160     goto cleanupAndFail;
161 }
162
163 } else {
164     /** Compress data.
165     **
166     ...
171     **/
172
173     unsigned int odatalen; /* maybe not the same size as outdatalen */
174     int blockSize100k = 9;
175
176     /* Get compression block size if present. */
177     if (cd_nelmts > 0) {
178         blockSize100k = cd_values[0];
179         if (blockSize100k < 1 || blockSize100k > 9) {
180             fprintf(stderr, "invalid compression block size: %d\n", blockSize100k);
181             goto cleanupAndFail;
182         }
183     }
184
185     /* Prepare the output buffer. */
186     outbuflen = nbytes + nbytes / 100 + 600; /* worst case (bzip2 docs) */
187     outbuf = malloc(outbuflen);
188     if (outbuf == NULL) {
189         fprintf(stderr, "memory allocation failed for bzip2 compression\n");
190         goto cleanupAndFail;
191     }
192
193     /* Compress data. */
194     odatalen = outbuflen;
195     ret = BZ2_bzBuffToBuffCompress(outbuf, &odatalen, *buf, nbytes,
196                                   blockSize100k, 0, 0);
197     outdatalen = odatalen;
198     if (ret != BZ_OK) {
199         fprintf(stderr, "bzip2 compression failed with error %d\n", ret);
200         goto cleanupAndFail;

```



```
201     }
202 }
203
204 /* Always replace the input buffer with the output buffer. */
205 free(*buf);
206 *buf = outbuf;
207 *buf_size = outbuflen;
208 return outdatalen;
209
210 cleanupAndFail:
211 if (outbuf)
212     free(outbuf);
213 return 0;
214 }
```

10 Example of Makefile

This is an example for Makefile. Building shared libraries may be tricky. Please consult system documentation for the flags on your system. You may also want to consider using `libtool` as we do in the `bzip2` plugin example.

```

1  CFLAGS = -fPIC -g
2  CC = gcc
3  HDF5_INSTALL = /mnt/scr1/_tmp/_build/hdf5/
4  BZIP2_INSTALL = /mnt/scr1/bzip2-1.0.5
5  MAJOR = 0
6  MINOR = 1
7  NAME1 = H5Zbzip2
8  VERSION = $(MAJOR).$(MINOR)
9
10 # Include files in hdf5/src build/src directories for hdf5.h and
    H5pubconf.h
11 INCLUDES = -I./ -I$(HDF5_INSTALL)/include -I$(BZIP2_INSTALL)
12
13 lib: lib$(NAME1).so.$(VERSION)
14
15 $(NAME1).o: $(NAME1).c
16     $(CC) $(CFLAGS) $(INCLUDES) -c $(NAME1).c
17
18 lib$(NAME1).so.$(VERSION): $(NAME1).o
19     $(CC) -shared -Wl,-soname,lib$(NAME1).so.$(MAJOR) $^ -o $@ -
L/$(BZIP2_INSTALL) -lbz2
20
21 clean:
22     $(RM) *.o *.so*
23
24 distclean:
25     $(RM) *.o *.so*
```

11 Revision History

<i>March 31, 2013</i>	Version 1 shared with developers who worked on the project.
<i>April 9, 2013</i>	Version 2 shared with The HDF Group developers.
<i>April 16, 2013</i>	Version 3: accepted changes from Mark

12 References

1. The HDF Group. "HDF5 Documentation," <http://www.hdfgroup.org/HDF5/doc/doc-info.html> (November 15, 2012).