

User Guide for Developing a Virtual Object Layer Plugin

Mohamad Chaarawi

12th September 2014

Contents

1	Introduction	1
2	Creating a VOL Plugin	1
2.1	Mapping the API to the Callbacks	3
2.2	The File Function Callbacks	5
2.3	The Group Function Callbacks	10
2.4	The Dataset Function Callbacks	13
2.5	The Attribute Function Callbacks	18
2.6	The Named Datatype Function Callbacks	23
2.7	The Object Function Callbacks	27
2.8	The Link Function Callbacks	31
2.9	The Asynchronous Function Callbacks	37
2.10	The Optional Generic Callback	38
3	New VOL API Routines	38
4	Creating and Using an Internal Plugin	43
4.1	Implementing an Internal Plugin	44
4.2	Using an Internal Plugin	45
5	Creating and Using an External Plugin	46
5.1	New API Routines for External Plugins	46
5.2	Using an External Plugin	46
6	Interchanging and Stacking VOL Plugins	47
6.1	Stacking Plugins on Top of Each Other	47
6.2	Mirroring Plugins	47
6.3	Implementing Stacked and Mirrored Plugins	48

1 Introduction

The Virtual Object Layer (VOL) is an abstraction layer in the HDF5 library that intercepts all API calls that could potentially access objects in an HDF5 container and forwards those calls to plugin "object drivers". The plugins could store the objects in variety of ways. A plugin could, for example, have objects be distributed remotely over different platforms, provide a raw mapping of the model to the file system, or even store the data in other file formats (like native netCDF or HDF4 format). The user still gets the same data model where access is done to a single HDF5 "container"; however the plugin object driver translates from what the user sees to how the data is actually stored. Having this abstraction layer maintains the object model of HDF5 and would allow HDF5 developers or users to write their own plugins for accessing HDF5 data.

This user guide is for developers interested in developing a VOL plugin for the HDF5 library. The document is meant to be used in conjunction with the HDF5 reference manual. It is assumed that the reader has good knowledge of the VOL architecture obtained by reading the VOL architectural design document ?? MSC-ref. The document will cover the steps needed to create external and internal VOL plugins. Both ways have a lot of common steps and rules that will be covered first.

2 Creating a VOL Plugin

Each VOL plugin should be of type `H5VL_class_t` that is defined as:

```
/* Class information for each VOL driver */
typedef struct H5VL_class_t {
    H5VL_class_value_t value;
    const char *name;
    herr_t (*initialize)(hid_t vpl_id);
    herr_t (*terminate)(hid_t vtpl_id);
    size_t fapl_size;
    void * (*fapl_copy)(const void *info);
    herr_t (*fapl_free)(void *info);

    /* Data Model */
    H5VL_attr_class_t      attr_cls;
    H5VL_dataset_class_t  dataset_cls;
    H5VL_datatype_class_t datatype_cls;
    H5VL_file_class_t      file_cls;
    H5VL_group_class_t     group_cls;
    H5VL_link_class_t      link_cls;
    H5VL_object_class_t    object_cls;

    /* Services */
    H5VL_async_class_t      async_cls;
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
        arguments);
} H5VL_class_t;
```

The **value** field is a unique integer enum identifier that should be greater than 128 for external plugins and smaller than 128 for internal plugins. Setting it in the VOL structure is required.

The **name** field is a string that uniquely identifies the VOL plugin name. Setting it in the VOL structure is required. No two plugins with the same name can be registered with the library at the same time.

The **initialize** field is a function pointer to a routine that a plugin implements to set up or initialize access to the plugin. Implementing this function by the plugin is not required, since some plugins do not require any set up to start accessing the plugin. In that case, the value of the function pointer should be set to NULL. Plugin specific variables that are required to be passed from users should be passed through the VOL initialize property list. Generic properties can be added to this property class for user-defined plugins that can not modify the HDF5 library to add internal properties. For more information consult the property list reference manual pages.

The **terminate** field is a function pointer to a routine that a plugin implements to terminate or finalize access to the plugin. Implementing this function by the plugin is not required, since some plugins do not require any termination phase to the plugin. In that case, the value of the function pointer should be set to NULL. Plugin specific variables that are required to be passed from users should be passed through the VOL terminate property list. Generic properties can be added to this property class for user-defined plugins that can not modify the HDF5 library to add internal properties. For more information consult the property list reference manual pages.

The **fapl_size** field indicates the size required to store the info data that the plugin needs. That info data is passed when the plugin is selected for usage with the file access property list (fapl) function. It might be that the plugin defined does not require any information from the user, which means the size in this field will be zero. More information about the info data and the fapl selection routines follow later.

The **fapl_copy** field is a function pointer that is called when the plugin is selected with the fapl function. It allows the plugin to make a copy of the info data since the user might free it when closing the fapl. It is required if there is info data needed by the plugin.

The **fapl_free** field is a function pointer that is called to free the info data when the fapl close routine is called. It is required if there is info data needed by the plugin.

The rest of the fields in the **H5VL_class_t** struct are "subclasses" that define all the VOL function callbacks that are mapped to from the HDF5 API layer. Those subclasses are categorized into two categories, Data Model and Services. Data Model classes are those that provide functionality for accessing an HDF5 container and objects in that container as defined by the HDF5 data model. Service classes are those that provide services for users that are not related to the data model specifically. Asynchronous operations, for example, are a service that most plugins can implement, so we add a class for it in the VOL structure. If a service becomes generic enough and common between many plugins, a class

for it should be added in the VOL structure. However, many plugins can/will provide services that are not shared by other plugins. A good way to support these services is through an optional callback in the VOL structure which can be a hook from the API to the plugin that provides those services, passing any necessary arguments needed without the HDF5 library having to worry about supporting that service. A similar API operation to allow users to use that service will be added. This API call would be similar to an “ioctl” call where any kind of operation can be supported and passed down to the plugin that has enough knowledge from the user to interpret the type of the operation. All classes and their defined callbacks will be detailed in the following subsections.

2.1 Mapping the API to the Callbacks

The callback interface defined for the VOL has to be general enough to handle all the HDF5 API operations that would access the file. Furthermore it has to capture future additions to the HDF5 library with little to no changes to the callback interface. Changing the interface often whenever new features are added would be discouraging to plugin developers since that would mean reworking their VOL plugin structure. To remedy this issue, every callback will contain two parameters:

- A data transfer property list (DXPL) which allows that API to put some properties on for the plugins to retrieve if they have to for particular operations, without having to add arguments to the VOL callback function.
- A pointer to a request (`void **req`) to handle asynchronous operations if the HDF5 library adds support for them in future releases (beyond the 1.8 series). That pointer is set by the VOL plugin to a request object it creates to manage progress on that asynchronous operation. If the `req` is `NULL`, that means that the API operation is blocking and so the plugin would not execute the operation asynchronously. If the plugin does not support asynchronous operations, it needs not to worry about this field and leaves it unset.

In order to keep the number of the VOL object classes and callbacks concise and readable, it was decided to not have a one-to-one mapping between API operation and callbacks. Furthermore, to keep the callbacks themselves short and not cluttered with a lot of parameters, some of the parameters are passed in as properties in property lists included with the callback. The value of those properties can be retrieved by calling the public routine (or its private version if this is an internal plugin):

```
herr_t H5Pget(hid_t plist_id, const char *property_name, void *value);
```

The property names and value types will be detailed when describing each callback in their respective sections.

The HDF5 library provides several routines to access an object in the container. For example to open an attribute on a group object, the user could use

H5Aopen() and pass the group identifier directly where the attribute needs to be opened. Alternatively, the user could use H5Aopen_by_name() or H5Aopen_by_idx() to open the attribute, which provides a more flexible way of locating the attribute, whether by a starting object location and a path or an index type and traversal order. All those types of accesses usually map to one VOL callback with a parameter that indicates the access type. In the example of opening an attribute, the three API open routine will map to the same VOL open callback but with a different location parameter. The same applies to all types of routines that have multiple types of accesses. The location parameter is a structure defined as follows:

```

/*
 * Structure to hold parameters for object locations.
 * either: BY_SELF, BY_NAME, BY_IDX, BY_ADDR, BY_REF
 */

typedef struct H5VL_loc_params_t {
    H5I_type_t obj_type; /* The object type of the location object */
    H5VL_loc_type_t type; /* The location type */
    union { /* parameters of the location */
        struct H5VL_loc_by_name loc_by_name;
        struct H5VL_loc_by_idx loc_by_idx;
        struct H5VL_loc_by_addr loc_by_addr;
        struct H5VL_loc_by_ref loc_by_ref;
    } loc_data;
} H5VL_loc_params_t

/*
 * Types for different ways that objects are located in an
 * HDF5 container.
 */
typedef enum H5VL_loc_type_t {
    /* starting location is the target object*/
    H5VL_OBJECT_BY_SELF,

    /* location defined by object and path in H5VL_loc_by_name */
    H5VL_OBJECT_BY_NAME,

    /* location defined by object, path, and index in H5VL_loc_by_idx */
    H5VL_OBJECT_BY_IDX,

    /* location defined by physical address in H5VL_loc_by_addr */
    H5VL_OBJECT_BY_ADDR,

    /* NOT USED */
    H5VL_OBJECT_BY_REF
} H5VL_loc_type_t;

struct H5VL_loc_by_name {
    const char *name; /* The path relative to the starting location */
    hid_t plist_id; /* The link access property list */
};

```

```

struct H5VL_loc_by_idx {
    const char *name; /* The path relative to the starting location */
    H5_index_t idx_type; /* Type of index */
    H5_iter_order_t order; /* Index traversal order */
    hsize_t n; /* position in index */
    hid_t plist_id; /* The link access property list */
};

struct H5VL_loc_by_addr {
    haddr_t addr; /* physical address of location */
};

/* Not used for now */
struct H5VL_loc_by_ref {
    H5R_type_t ref_type;
    const void *_ref;
    hid_t plist_id;
};

```

Other API routines that would make a one-to-one mapping difficult are:

- the **Get** operations that retrieve something from an object; for example a property list or a datatype of a dataset, etc...
- API routines specific to each class of callbacks(attributes, files, etc...). Having a callback for each one would clutter the VOL callback structure and reduce user friendliness.
- API routines that make sense to only implement in the native HDF5 file format and would not be possible to implement in other plugins. Adding a callback for those routines would be confusing to plugin developers.

To handle that large set of API routines, each class in the Data Model category has three generic callbacks, **get**, **specific**, and **optional** to handle the three set of API operations outline above respectively. The callbacks will have a **va_list** argument to handle the different set of parameters that could be passed in. The **get** and **specific** callbacks also have an **op_type** that contain an **enum** of the type of operation that is requested to be performed. Using that type, the **va_list** argument can be parsed. The **optional** callback is a free for all callback where anything from the API layer is passed in directly. This callback is used to support plugin specific operations in the API that other plugins should or would not know about. More information about types and the arguments for each type will be detailed in the corresponding class arguments.

2.2 The File Function Callbacks

The file API routines (H5F) allow HDF5 users to create and manage HDF5 containers. All the H5F API routines that modify the HDF5 container map to one of the file callback routines in this class that the plugin needs to implement:

```

typedef struct H5VL_file_class_t {

```

```

void *(*create)(const char *name, unsigned flags, hid_t fcpl_id,
               hid_t fapl_id, hid_t dxpl_id, void **req);
void *(*open)(const char *name, unsigned flags, hid_t fapl_id, hid_t
              dxpl_id, void **req);
herr_t (*get)(void *obj, H5VL_file_get_t get_type, hid_t dxpl_id,
              void **req, va_list arguments);
herr_t (*specific)(void *obj, H5VL_file_specific_t specific_type,
                  hid_t dxpl_id, void **req, va_list arguments);
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
                  arguments);
herr_t (*close)(void *file, hid_t dxpl_id, void **req);
} H5VL_file_class_t;

```

The `create` callback in the file class should create a container and returns a pointer to the file structure created by the plugin containing information to access the container in future calls.

Signature:

```

void *(*create)(const char *name, unsigned flags, hid_t fcpl_id,
               hid_t fapl_id, hid_t dxpl_id, void **req);

```

Arguments:

name (IN): The name of the container to be created.
flags (IN): The creation flags of the container.
fcpl_id (IN): The file creation property list.
fapl_id (IN): The file access property list.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `open` callback in the file class should open a container and returns a pointer to the file structure created by the plugin containing information to access the container in future calls.

Signature:

```

void *(*open)(const char *name, unsigned flags, hid_t fapl_id, hid_t
              dxpl_id, void **req);

```

Arguments:

name (IN): The name of the container to open.
flags (IN): The open flags of the container.
fapl_id (IN): The file access property list.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the file class should retrieve information about the container as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *obj, H5VL_file_get_t get_type, hid_t dxpl_id,  
              void **req, va_list arguments);
```

The `get_type` argument is an enum:

```
/* types for all file get API routines */  
typedef enum H5VL_file_get_t {  
    H5VL_FILE_GET_FAPL, /* file access property list */  
    H5VL_FILE_GET_FCPL, /* file creation property list */  
    H5VL_FILE_GET_INTENT, /* file intent */  
    H5VL_FILE_GET_NAME, /* file name */  
    H5VL_FILE_GET_OBJ_COUNT, /* object count in file */  
    H5VL_FILE_GET_OBJ_IDS, /* object ids in file */  
    H5VL_OBJECT_GET_FILE /* retrieve or resurrect file of object */  
} H5VL_file_get_t;
```

Arguments:

<code>obj</code>	(IN): The container or object where information needs to be retrieved from.
<code>get_type</code>	(IN): The type of the information to retrieve.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `get_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_FILE_GET_FCPL`, to retrieve the file creation property list:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the file creation property list.
- `H5VL_FILE_GET_FAPL`, to retrieve the file access property list:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the file access property list.
- `H5VL_FILE_GET_OBJ_COUNT`:, to retrieve the object count in the container:
 1. `unsigned types` (IN): type of objects to look for.
 2. `ssize_t *ret` (OUT): buffer for the object count.
- `H5VL_FILE_GET_OBJ_IDS`:, to retrieve object identifiers in the container:
 1. `unsigned types` (IN): type of objects to look for.
 2. `size_t max_objs` (IN): maximum number of objects to open.
 3. `hid_t *oid_list` (OUT): buffer for the object identifiers.
 4. `ssize_t *ret` (OUT): buffer for the object count.

- `H5VL_FILE_GET_INTENT`, get access intent of the container:
 1. `unsigned *ret` (OUT): buffer for the intent value.
- `H5VL_FILE_GET_NAME`, get container name an object belongs to:
 1. `H5I_type_t type` (IN): the object type in `obj`.
 2. `size_t size` (IN): size of the buffer for the file name.
 3. `char *name` (OUT): buffer for the file name.
 4. `ssize_t *ret` (OUT): buffer for the entire size of the file name.
- `H5VL_OBJECT_GET_FILE`, get the container that the object belongs to:
 1. `H5I_type_t type` (IN): the object type in `obj`.
 2. `void **ret` (OUT): pointer to the file structure returned by the plugin.

The `specific` callback in the file class implements specific operations on HDF5 files as specified in the `specific_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*specific)(void *obj, H5VL_file_specific_t specific_type,
                  hid_t dxpl_id, void **req, va_list arguments);
```

The `specific_type` argument is an `enum`:

```
/* types for file SPECIFIC callback */
typedef enum H5VL_file_specific_t {
    H5VL_FILE_FLUSH,                /* Flush file */
    H5VL_FILE_IS_ACCESSIBLE,        /* Check if a file is accessible */
    H5VL_FILE_MOUNT,                /* Mount a file */
    H5VL_FILE_UNMOUNT               /* Un-Mount a file */
} H5VL_file_specific_t;
```

Arguments:

<code>obj</code>	(IN): The container or object where the operation needs to happen.
<code>specific_type</code>	(IN): The type of the operation.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `specific_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_FILE_FLUSH`, flushes all buffers associated with the container to disk:

1. `H5I_type_t obj_type` (IN): The object type of `obj` on which the flush operation was called.
 2. `H5F_scope_t scope` (IN): The scope of the flushing action.
- `H5VL_FILE_IS_ACCESSIBLE`, checks if a container is accessible using a specific file access property list:
 1. `hid_t *fapl_id` (IN): file access property list.
 2. `char *name` (IN): name of the container to check.
 3. `htri_t *result` (OUT): buffer for the result; 0 if no, 1 if yes.
 - `H5VL_FILE_MOUNT`, Mounts a file on the location object:
 1. `H5I_type_t type` (IN): the object type in `obj`.
 2. `char *name` (IN): name of the group onto which the file specified by `file` is to be mounted.
 3. `void *file` (IN): child file to be mounted.
 4. `hid_t *fmpl_id` (IN): file mount property list.
 - `H5VL_FILE_UNMOUNT`, un-mounts a file from the location object:
 1. `H5I_type_t type` (IN): the object type in `obj`.
 2. `char *name` (IN): name of the mount point.

The optional callback in the file class implements plugin specific operations on an HDF5 container. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
arguments);
```

Arguments:

<code>obj</code>	(IN): The container or object where the operation needs to happen.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema. For example, in the native plugin, the first argument in the `va_list` is the optional argument type enum defined as follows:

```
/* types for all file optional operations */
typedef enum H5VL_file_optional_t {
    H5VL_FILE_CLEAR_ELINK_CACHE, /* Clear external link cache */
    H5VL_FILE_GET_FILE_IMAGE,   /* file image */
}
```

```

H5VL_FILE_GET_FREE_SECTIONS, /* file free selections      */
H5VL_FILE_GET_FREE_SPACE,   /* file freespace      */
H5VL_FILE_GET_INFO,         /* file info           */
H5VL_FILE_GET_MDC_CONF,     /* file metadata cache configuration */
H5VL_FILE_GET_MDC_HR,       /* file metadata cache hit rate */
H5VL_FILE_GET_MDC_SIZE,     /* file metadata cache size */
H5VL_FILE_GET_SIZE,         /* file size           */
H5VL_FILE_GET_VFD_HANDLE,   /* file VFD handle     */
H5VL_FILE_REOPEN,           /* reopen the file      */
H5VL_FILE_RESET_MDC_HIT_RATE, /* get metadata cache hit rate */
H5VL_FILE_SET_MDC_CONFIG    /* set metadata cache configuration */
} H5VL_file_optional_t;

```

The remaining arguments are parsed depending on the optional operation type.

The `close` callback in the file class should terminate access to the file object and free all resources it was consuming, and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*close)(void *file, hid_t dxpl_id, void **req);
```

Arguments:

`file` (IN): Pointer to the file.
`dxpl_id` (IN): The data transfer property list.
`req` (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

2.3 The Group Function Callbacks

The group API routines (H5G) allow HDF5 users to create and manage HDF5 groups. All the H5G API routines that modify the HDF5 container map to one of the group callback routines in this class that the plugin needs to implement:

```

typedef struct H5VL_group_class_t {
    void *(*create)(void *obj, H5VL_loc_params_t loc_params, const char
        *name, hid_t gcpl_id, hid_t gapl_id, hid_t dxpl_id, void **req);
    void *(*open)(void *obj, H5VL_loc_params_t loc_params, const char
        *name, hid_t gapl_id, hid_t dxpl_id, void **req);
    herr_t (*get)(void *obj, H5VL_group_get_t get_type, hid_t dxpl_id,
        void **req, va_list arguments);
    herr_t (*specific)(void *obj, H5VL_group_specific_t specific_type,
        hid_t dxpl_id, void **req, va_list arguments);
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
        arguments);
    herr_t (*close)(void *grp, hid_t dxpl_id, void **req);
} H5VL_group_class_t;

```

The `create` callback in the group class should create a group object in the container of the location object and returns a pointer to the group structure containing information to access the group in future calls.

Signature:

```
void *(*create)(void *obj, H5VL_loc_params_t loc_params, const char
                *name, hid_t gcpl_id, hid_t gapl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the group needs to be created or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_SELF</code> in this callback.
<code>name</code>	(IN): The name of the group to be created.
<code>dcpl_id</code>	(IN): The group creation property list. It contains all the group creation properties in addition to the link creation property list of the create operation (an <code>hid_t</code>) that can be retrieved with the property <code>H5VL_GRP_LCPL_ID</code> .
<code>gapl_id</code>	(IN): The group access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `open` callback in the group class should open a group object in the container of the location object and returns a pointer to the group structure containing information to access the group in future calls.

Signature:

```
void *(*open)(void *obj, H5VL_loc_params_t loc_params,
              const char*name, hid_t gapl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the group needs to be opened or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_SELF</code> in this callback.
<code>name</code>	(IN): The name of the group to be opened.
<code>dapl_id</code>	(IN): The group access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the group class should retrieve information about the group as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *obj, H5VL_group_get_t get_type, hid_t dxpl_id,
              void **req, va_list arguments);
```

The `get_type` argument is an `enum`:

```
/* types for all group get API routines */
typedef enum H5VL_group_get_t {
    H5VL_GROUP_GET_GCPL,    /*group creation property list */
    H5VL_GROUP_GET_INFO    /*group info */
} H5VL_group_get_t;
```

Arguments:

- | | |
|------------------------|--|
| <code>obj</code> | (IN): The group object where information needs to be retrieved from. |
| <code>get_type</code> | (IN): The type of the information to retrieve. |
| <code>dxpl_id</code> | (IN): The data transfer property list. |
| <code>req</code> | (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin. |
| <code>arguments</code> | (IN/OUT): argument list containing parameters and output pointers for the get operation. |

The `arguments` argument contains a variable list of arguments depending on the `get_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_GROUP_GET_GCPL`, to retrieve the group creation property list of the group specified in `obj`:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the group creation property list.
- `H5VL_GROUP_GET_INFO`, to retrieve the attribute info:
 1. `H5VL_loc_params_t loc_params` (IN): The location parameters explained in section 2.1.
 2. `H5G_info_t *ginfo` (OUT): info structure to fill the group info in.

The `specific` callback in the group class currently does not require any functionality to be implemented and should be left as `NULL`.

The `optional` callback in the group class implements plugin specific operations on an HDF5 group. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
    arguments);
```

Arguments:

- | | |
|------------------------|--|
| <code>obj</code> | (IN): The container or object where the operation needs to happen. |
| <code>dxpl_id</code> | (IN): The data transfer property list. |
| <code>req</code> | (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin. |
| <code>arguments</code> | (IN/OUT): argument list containing parameters and output pointers for the get operation. |

Each plugin should be able to parse the `va_list` arguments if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

The `close` callback in the group class should terminate access to the group object and free all resources it was consuming, and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*close)(void *group, hid_t dxpl_id, void **req);
```

Arguments:

`group` (IN): Pointer to the group object.
`dxpl_id` (IN): The data transfer property list.
`req` (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

2.4 The Dataset Function Callbacks

The dataset API routines (H5D) allow HDF5 users to create and manage HDF5 datasets. All the H5D API routines that modify the HDF5 container map to one of the dataset callback routines in this class that the plugin needs to implement:

```
typedef struct H5VL_dataset_class_t {
    void *(*create)(void *obj, H5VL_loc_params_t loc_params, const char
        *name, hid_t dcpl_id, hid_t dapl_id, hid_t dxpl_id, void **req);
    void *(*open)(void *obj, H5VL_loc_params_t loc_params, const char
        *name, hid_t dapl_id, hid_t dxpl_id, void **req);
    herr_t (*read)(void *dset, hid_t mem_type_id, hid_t mem_space_id,
        hid_t file_space_id, hid_t xfer_plist_id, void *buf, void
        **req);
    herr_t (*write)(void *dset, hid_t mem_type_id, hid_t mem_space_id,
        hid_t file_space_id, hid_t xfer_plist_id, const void *buf, void
        **req);
    herr_t (*get)(void *obj, H5VL_dataset_get_t get_type, hid_t dxpl_id,
        void **req, va_list arguments);
    herr_t (*specific)(void *obj, H5VL_dataset_specific_t specific_type,
        hid_t dxpl_id, void **req, va_list arguments);
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
        arguments);
    herr_t (*close)(void *dset, hid_t dxpl_id, void **req);
} H5VL_dataset_class_t;
```

The `create` callback in the dataset class should create a dataset object in the container of the location object and returns a pointer to the dataset structure containing information to access the dataset in future calls.

Signature:

```
void *(*create)(void *obj, H5VL_loc_params_t loc_params, const char
                 *name, hid_t dcpl_id, hid_t dapl_id, hid_t dxpl_id, void **req);
```

Arguments:

obj	(IN): Pointer to an object where the dataset needs to be created or where the look-up of the target object needs to start.
loc_params	(IN): The location parameters as explained in section 2.1. The type can be only H5VL_OBJECT_BY_SELF in this callback.
name	(IN): The name of the dataset to be created.
dcpl_id	(IN): The dataset creation property list. It contains all the dataset creation properties in addition to the dataset datatype (an hid_t), dataspace (an hid_t), and the link creation property list of the create operation (an hid_t) that can be retrieved with the properties, H5VL_DSET_TYPE_ID, H5VL_DSET_SPACE_ID, and H5VL_DSET_LCPL_ID respectively.
dapl_id	(IN): The dataset access property list.
dxpl_id	(IN): The data transfer property list.
req	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `open` callback in the dataset class should open a dataset object in the container of the location object and returns a pointer to the dataset structure containing information to access the dataset in future calls.

Signature:

```
void *(*open)(void *obj, H5VL_loc_params_t loc_params, const char
              *name, hid_t dapl_id, hid_t dxpl_id, void **req);
```

Arguments:

obj	(IN): Pointer to an object where the dataset needs to be opened or where the look-up of the target object needs to start.
loc_params	(IN): The location parameters as explained in section 2.1. The type can be only H5VL_OBJECT_BY_SELF in this callback.
name	(IN): The name of the dataset to be opened.
dapl_id	(IN): The dataset access property list.
dxpl_id	(IN): The data transfer property list.
req	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `read` callback in the dataset class should read data from the dataset object and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*read)(void *dset, hid_t mem_type_id, hid_t mem_space_id,
               hid_t file_space_id, hid_t dxpl_id, void *buf, void **req);
```

Arguments:

<code>dset</code>	(IN): Pointer to the dataset object.
<code>mem_type_id</code>	(IN): The memory datatype of the data.
<code>mem_space_id</code>	(IN): The memory dataspace selection.
<code>file_space_id</code>	(IN): The file dataspace selection.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>buf</code>	(OUT): Data buffer to be read into.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `write` callback in the dataset class should write data to the dataset object and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*write)(void *dset, hid_t mem_type_id, hid_t mem_space_id,
               hid_t file_space_id, hid_t dxpl_id, const void * buf, void
               **req);
```

Arguments:

<code>dset</code>	(IN): Pointer to the dataset object.
<code>mem_type_id</code>	(IN): The memory datatype of the data.
<code>mem_space_id</code>	(IN): The memory dataspace selection.
<code>file_space_id</code>	(IN): The file dataspace selection.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>buf</code>	(IN): Data buffer to be written from.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the dataset class should retrieve information about the dataset as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *dset, H5VL_dataset_get_t get_type,
              hid_t dxpl_id, void **req, va_list arguments);
```

The `get_type` argument is an enum:

```
/* types for dataset GET callback */
typedef enum H5VL_dataset_get_t {
    H5VL_DATASET_GET_DAPL,           /* access property list */
    H5VL_DATASET_GET_DCPL,          /* creation property list */
    H5VL_DATASET_GET_OFFSET,        /* offset */
    H5VL_DATASET_GET_SPACE,         /* dataspace */
    H5VL_DATASET_GET_SPACE_STATUS,  /* space status */
    H5VL_DATASET_GET_STORAGE_SIZE,  /* storage size */
    H5VL_DATASET_GET_TYPE           /* datatype */
} H5VL_dataset_get_t;
```

Arguments:

dset	(IN): The dataset object where information needs to be retrieved from.
get_type	(IN): The type of the information to retrieve.
dxpl_id	(IN): The data transfer property list.
req	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
arguments	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The **arguments** argument contains a variable list of arguments depending on the **get_type** parameter. The following list shows the argument list, in-order, for each type:

- **H5VL_DATASET_GET_DAPL**, to retrieve the dataset access property list of the dataset specified in **obj**:
 1. **hid_t *ret_id** (OUT): buffer for the identifier of the dataset access property list.
- **H5VL_DATASET_GET_DCPL**, to retrieve the dataset creation property list of the dataset specified in **obj**:
 1. **hid_t *ret_id** (OUT): buffer for the identifier of the dataset creation property list.
- **H5VL_DATASET_GET_OFFSET**, to retrieve the offset of the dataset specified in **obj** in the container:
 1. **haddr_t *ret** (OUT): buffer for the offset of the dataset in the container.
- **H5VL_DATASET_GET_SPACE**, to retrieve the dataspace of the dataset specified in **obj**:
 1. **hid_t *ret_id** (OUT): buffer for the identifier of the dataset dataspace.
- **H5VL_DATASET_GET_SPACE_STATUS**, to retrieve the information whether space has been allocated for the dataset:
 1. **H5D_space_status_t *allocation** (OUT): buffer for the space status.
- **H5VL_DATASET_GET_STORAGE_SIZE**, to retrieve the storage size of the dataset specified in **obj**:
 1. **hsize_t *ret** (OUT): buffer for the storage size of the dataset in the container.
- **H5VL_DATASET_GET_TYPE**, to retrieve the datatype of the dataset specified in **obj**:
 1. **hid_t *ret_id** (OUT): buffer for the identifier of the dataset datatype.

The **specific** callback in the dataset class implements specific operations on HDF5 datasets as specified in the **specific_type** parameter. It returns an **herr_t** indicating success or failure.

Signature:

```
herr_t (*specific)(void *obj, H5VL_file_specific_t specific_type,  
                    hid_t dxpl_id, void **req, va_list arguments);
```

The `specific_type` argument is an enum:

```
/* types for dataset SPECIFIC callback */  
typedef enum H5VL_dataset_specific_t {  
    H5VL_DATASET_SET_EXTENT /* H5Dset_extent */  
} H5VL_dataset_specific_t;
```

Arguments:

<code>obj</code>	(IN): The dset where the operation needs to happen.
<code>specific_type</code>	(IN): The type of the operation.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `specific_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_DATASET_SET_EXTENT`, changes the extent of the dataset dimensions:
 1. `hsize_t size` (IN): Array containing the new magnitude of each dimension of the dataset.

The `optional` callback in the dataset class implements plugin specific operations on an HDF5 dataset. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list  
                    arguments);
```

Arguments:

<code>obj</code>	(IN): The container or object where the operation needs to happen.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

The `close` callback in the dataset class should terminate access to the dataset object and free all resources it was consuming, and returns an `herr_t` indicating

success or failure.

Signature:

```
herr_t (*close)(void *dset, hid_t dxpl_id, void **req);
```

Arguments:

dset (IN): Pointer to the dataset object.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

2.5 The Attribute Function Callbacks

The attribute API routines (H5A) allow HDF5 users to create and manage HDF5 attributes. All the H5A API routines that modify the HDF5 container map to one of the attribute callback routines in this class that the plugin needs to implement:

```
typedef struct H5VL_attr_class_t {  
    void *(*create)(void *obj, H5VL_loc_params_t loc_params, const char  
        *attr_name, hid_t acpl_id, hid_t aapl_id, hid_t dxpl_id, void  
        **req);  
    void *(*open)(void *obj, H5VL_loc_params_t loc_params, const char  
        *attr_name, hid_t aapl_id, hid_t dxpl_id, void **req);  
    herr_t (*read)(void *attr, hid_t mem_type_id, void *buf, hid_t  
        dxpl_id, void **req);  
    herr_t (*write)(void *attr, hid_t mem_type_id, const void *buf,  
        hid_t dxpl_id, void **req);  
    herr_t (*get)(void *obj, H5VL_attr_get_t get_type, hid_t dxpl_id,  
        void **req, va_list arguments);  
    herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,  
        H5VL_attr_specific_t specific_type, hid_t dxpl_id, void **req,  
        va_list arguments);  
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list  
        arguments);  
    herr_t (*close)(void *attr, hid_t dxpl_id, void **req);  
} H5VL_attr_class_t;
```

The `create` callback in the attribute class should create an attribute object in the container of the location object and returns a pointer to the attribute structure containing information to access the attribute in future calls.

Signature:

```
void *(*create)(void *obj, H5VL_loc_params_t loc_params,  
    const char *attr_name, hid_t acpl_id, hid_t aapl_id,  
    hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the attribute needs to be created or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>attr_name</code>	(IN): The name of the attribute to be created.
<code>acpl_id</code>	(IN): The attribute creation property list. It contains all the attribute creation properties in addition to the attribute datatype (an <code>hid_t</code>) and dataspace (an <code>hid_t</code>) that can be retrieved with the properties, <code>H5VL_ATTR_TYPE_ID</code> and <code>H5VL_ATTR_SPACE_ID</code> .
<code>aapl_id</code>	(IN): The attribute access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `open` callback in the attribute class should open an attribute object in the container of the location object and returns a pointer to the attribute structure containing information to access the attribute in future calls.

Signature:

```
void *(*open)(void *obj, H5VL_loc_params_t loc_params,
              const char *attr_name, hid_t aapl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the attribute needs to be opened or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>attr_name</code>	(IN): The name of the attribute to be opened.
<code>aapl_id</code>	(IN): The attribute access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `read` callback in the attribute class should read data from the attribute object and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*read)(void *attr, hid_t mem_type_id, void *buf,
              hid_t dxpl_id, void **req);
```

Arguments:

<code>attr</code>	(IN): Pointer to the attribute object.
<code>mem_type_id</code>	(IN): The memory datatype of the attribute.
<code>buf</code>	(OUT): Data buffer to be read into.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `write` callback in the attribute class should write data to the attribute object and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*write)(void *attr, hid_t mem_type_id, const void *buf,
                hid_t dxpl_id, void **req);
```

Arguments:

attr (IN): Pointer to the attribute object.
mem_type_id (IN): The memory datatype of the attribute.
buf (IN): Data buffer to be written.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the attribute class should retrieve information about the attribute as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *obj, H5VL_attr_get_t get_type, hid_t dxpl_id,
              void **req, va_list arguments);
```

The `get_type` argument is an enum:

```
/* types for attribute GET callback */
typedef enum H5VL_attr_get_t {
    H5VL_ATTR_GET_ACPL,          /* creation property list */
    H5VL_ATTR_GET_INFO,         /* info */
    H5VL_ATTR_GET_NAME,         /* access property list */
    H5VL_ATTR_GET_SPACE,        /* dataspace */
    H5VL_ATTR_GET_STORAGE_SIZE, /* storage size */
    H5VL_ATTR_GET_TYPE          /* datatype */
} H5VL_attr_get_t;
```

Arguments:

attr (IN): An attribute or location object where information needs to be retrieved from.
get_type (IN): The type of the information to retrieve.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
arguments (IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `get_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_ATTR.GET_ACPL`, to retrieve the attribute creation property list of the attribute specified in `obj`:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the attribute creation property list.

- `H5VL_ATTR.GET_INFO`, to retrieve the attribute info:
 1. `H5VL_loc_params_t loc_params` (IN): The location parameters explained in section 2.1.
 2. `H5A_info_t *ainfo` (OUT): info structure to fill the attribute info in.
- `H5VL_ATTR.GET_NAME`, to retrieve an attribute name on a particular object specified in `obj`:
 1. `H5VL_loc_params_t loc_params` (IN): The location parameters explained in section 2.1. The type could be either `H5VL_OBJECT_BY_SELF` meaning `obj` is the attribute, or `H5VL_OBJECT_BY_IDX` meaning the attribute to retrieve the name for should be looked up using the index information on the object in `obj` and the index information in `loc_params`.
 2. `size_t buf_size` (IN): the size of the buffer to store the name in.
 3. `void *buf` (OUT): Buffer to store the name in.
 4. `ssize_t *ret_val` (OUT): return the actual size needed to store the fill attribute name.
- `H5VL_ATTR.GET_SPACE`, to retrieve the dataspace of the attribute specified in `obj`:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the attribute dataspace.
- `H5VL_ATTR.GET_STORAGE_SIZE`, to retrieve the storage size of the attribute specified in `obj`:
 1. `hsize_t *ret` (OUT): buffer for the storage size of the attribute in the container.
- `H5VL_ATTR.GET_TYPE`, to retrieve the datatype of the attribute specified in `obj`:
 1. `hid_t *ret_id` (OUT): buffer for the identifier of the attribute datatype.

The `specific` callback in the attribute class implements specific operations on HDF5 attributes as specified in the `specific_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,
                  H5VL_attr_specific_t specific_type, hid_t dxpl_id, void
                  **req, va_list arguments);
```

The `specific_type` argument is an enum:

```
/* types for attribute SPECIFIC callback */
typedef enum H5VL_attr_specific_t {
    H5VL_ATTR_DELETE,                /* H5Adelete(_by_name/idx) */
```

```

H5VL_ATTR_EXISTS,          /* H5Aexists(_by_name) */
H5VL_ATTR_ITER,           /* H5Aiterate(_by_name) */
H5VL_ATTR_RENAME          /* H5Arename(_by_name) */
} H5VL_attr_specific_t;

```

Arguments:

<code>obj</code>	(IN): The location object where the operation needs to happen.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>specific_type</code>	(IN): The type of the operation.
<code>dexpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `specific_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_ATTR_DELETE`, to delete an attribute on an object:
 1. `char *attr_name` (IN): the name of the attribute to delete.
- `H5VL_ATTR_EXISTS`, to check if an attribute exists on a particular object specified in `obj`:
 1. `H5VL_loc_params_t loc_params` (IN): The location parameters explained in section 2.1.
 2. `char *attr_name` (IN): the attribute name to check.
 3. `htri_t *ret` (OUT): existence result, 0 if false, 1 if true.
- `H5VL_ATTR_ITER`, to iterate over all attributes of an object and call a user specified callback on each one:
 1. `H5_index_t idx_type` (IN): Type of index.
 2. `H5_iter_order_t order` (IN): Order in which to iterate over index.
 3. `hsize_t *idx` (IN/OUT): Initial and return offset within index.
 4. `H5A_operator2_t op` (IN): User-defined function to pass each attribute to.
 5. `void *op_data` (IN/OUT): User data to pass through to and to be returned by iterator operator function.
- `H5VL_ATTR_RENAME`, to rename an attribute on an object:
 1. `char *old_name` (IN): the original name of the attribute.

2. `char *new_name` (IN): the new name to assign for the attribute.

The `optional` callback in the attribute class implements plugin specific operations on an HDF5 attribute. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
arguments);
```

Arguments:

`obj` (IN): The container or object where the operation needs to happen.
`dxpl_id` (IN): The data transfer property list.
`req` (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
`arguments` (IN/OUT): argument list containing parameters and output pointers for the get operation.

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

The `close` callback in the attribute class should terminate access to the attribute object and free all resources it was consuming, and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*close)(void *attr, hid_t dxpl_id, void **req);
```

Arguments:

`attr` (IN): Pointer to the attribute object.
`dxpl_id` (IN): The data transfer property list.
`req` (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

2.6 The Named Datatype Function Callbacks

The HDF5 datatype routines (H5T) allow users to create and manage HDF5 datatypes. Those routines are divided into two categories. One that operates on all types of datatypes but do not modify the contents of the container (all in memory), and others that operate on named datatypes by accessing the container. When a user creates an HDF5 datatype, it is still an object in memory space (transient datatype) that has not been added to the HDF5 containers. Only when a user commits the HDF5 datatype, it becomes persistent in the container. Those are called named/committed datatypes. The transient H5T routines should work on named datatypes nevertheless.

All the H5T API routines that modify the HDF5 container map to one of the named datatype callback routines in this class that the plugin needs to

implement:

```
typedef struct H5VL_datatype_class_t {
    void *(*commit)(void *obj, H5VL_loc_params_t loc_params, const char
        *name, hid_t type_id, hid_t lcpl_id, hid_t tcpl_id, hid_t
        tapl_id, hid_t dxpl_id, void **req);
    void *(*open)(void *obj, H5VL_loc_params_t loc_params, const char *
        name, hid_t tapl_id, hid_t dxpl_id, void **req);
    herr_t (*get) (void *obj, H5VL_datatype_get_t get_type, hid_t
        dxpl_id, void **req, va_list arguments);
    herr_t (*specific)(void *obj, H5VL_datatype_specific_t
        specific_type, hid_t dxpl_id, void **req, va_list arguments);
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
        arguments);
    herr_t (*close) (void *dt, hid_t dxpl_id, void **req);
} H5VL_datatype_class_t;
```

The `commit` callback in the named datatype class should create a datatype object in the container of the location object and returns a pointer to the datatype structure containing information to access the datatype in future calls.

Signature:

```
void *(*commit)(void *obj, H5VL_loc_params_t loc_params,
    const char *name, hid_t type_id, hid_t lcpl_id, hid_t tcpl_id,
    hid_t tapl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the datatype needs to be committed or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1. In this call, the location type is always <code>H5VL_OBJECT_BY_SELF</code> .
<code>name</code>	(IN): The name of the datatype to be created.
<code>type_id</code>	(IN): The transient datatype identifier to be committed.
<code>lcpl_id</code>	(IN): The link creation property list.
<code>tcpl_id</code>	(IN): The datatype creation property list.
<code>tapl_id</code>	(IN): The datatype access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `open` callback in the named datatype class should open a previously committed datatype object in the container of the location object and returns a pointer to the datatype structure containing information to access the datatype in future calls.

Signature:

```
void *(*open) (void *obj, H5VL_loc_params_t loc_params,
    const char * name, hid_t tapl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to an object where the datatype needs to be opened or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1. In this call, the location type is always <code>H5VL_OBJECT_BY_SELF</code> .
<code>name</code>	(IN): The name of the datatype to be opened.
<code>tapl_id</code>	(IN): The datatype access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get_binary` callback in the named datatype class should serialize the original transient HDF5 datatype that was committed, or return the size that is required for it be serialized if the passed in buffer is `NULL`. The HDF5 library provides two functions to encode and decode datatypes in their transient form, `H5Tencode()` and `H5Tdecode()`. When a datatype is committed, the plugin is required to keep the serialized form of the transient datatype stored somewhere in the container (which is usually the case anyway when committing a named datatype), so it can be retrieved with this call. This is needed to generate the higher level HDF5 datatype identifier that allows all the H5T “transient” routines to work properly on the named datatype.

Signature:

```
ssize_t (*get_binary)(void *obj, unsigned char *buf, size_t size,
                     hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to the named datatype object.
<code>buf</code>	(OUT): Buffer to out the binary form of the datatype in.
<code>size</code>	(IN): The size of the buffer passed in (0 if <code>NULL</code>).
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the named datatype class should retrieve information about the named datatype as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get) (void *obj, H5VL_datatype_get_t get_type,
              hid_t dxpl_id, void **req, va_list arguments);
```

The `get_type` argument is an enum:

```
/* types for datatype GET callback */
typedef enum H5VL_datatype_get_t {
    H5VL_DATATYPE_GET_BINARY,          /* get serialized form of
    transient type */
    H5VL_DATATYPE_GET_TCPL             /* datatype creation property
    list */
} H5VL_datatype_get_t;
```

Arguments:

- obj** (IN): The named datatype to retrieve information from.
- get_type** (IN): The type of the information to retrieve.
- dxpl_id** (IN): The data transfer property list.
- req** (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
- arguments** (IN/OUT): argument list containing parameters and output pointers for the get operation.

The **arguments** argument contains a variable list of arguments depending on the **get_type** parameter. The following list shows the argument list, in-order, for each type:

- **H5VL_DATATYPE_GET_BINARY**, to retrieve the serialized original transient HDF5 datatype that was committed, or return the size that is required for it be serialized if the passed in buffer is **NULL**. The HDF5 library provides two functions to encode and decode datatypes in their transient form, **H5Tencode()** and **H5Tdecode()**. When a datatype is committed, the plugin is required to keep the serialized form of the transient datatype stored somewhere in the container (which is usually the case anyway when committing a named datatype), so it can be retrieved with this call. This is needed to generate the higher level HDF5 datatype identifier that allows all the H5T “transient” routines to work properly on the named datatype.
 1. **ssize_t *nalloc** (OUT): buffer to store the total size of the serialized datatype.
 2. **void *buf** (OUT): buffer to store the serialized datatype.
 3. **size_t size** (IN): size of the passed in buffer.
- **H5VL_DATATYPE_GET_TCPL**, to retrieve the datatype creation property list:
 1. **hid_t *ret_id** (OUT): buffer for the identifier of the type creation property list.

The **specific** callback in the datatype class currently does not require any functionality to be implemented and should be left as **NULL**.

The **optional** callback in the datatype class implements plugin specific operations on an HDF5 datatype. It returns an **herr_t** indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
                  arguments);
```

Arguments:

<code>obj</code>	(IN): The container or object where the operation needs to happen.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

The `close` callback in the named datatype class should terminate access to the datatype object and free all resources it was consuming, and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*close) (void *dt, hid_t dxpl_id, void **req);
```

Arguments:

<code>dt</code>	(IN): Pointer to the datatype object.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

2.7 The Object Function Callbacks

The object API routines (H5O) allow HDF5 users to manage HDF5 group, dataset, and named datatype objects. All the H5O API routines that modify the HDF5 container map to one of the object callback routines in this class that the plugin needs to implement:

```
typedef struct H5VL_object_class_t {
    void *(*open)(void *obj, H5VL_loc_params_t loc_params, H5I_type_t
        *opened_type, hid_t dxpl_id, void **req);
    herr_t (*copy)(void *src_obj, H5VL_loc_params_t loc_params1, const
        char *src_name, void *dst_obj, H5VL_loc_params_t loc_params2,
        const char *dst_name, hid_t ocpypl_id, hid_t lcpl_id, hid_t
        dxpl_id, void **req);
    herr_t (*get)(void *obj, H5VL_loc_params_t loc_params,
        H5VL_object_get_t get_type, hid_t dxpl_id, void **req, va_list
        arguments);
    herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,
        H5VL_object_specific_t specific_type, hid_t dxpl_id, void **req,
        va_list arguments);
    herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
        arguments);
} H5VL_object_class_t;
```

The `open` callback in the object class should open the object in the container of the location object and returns a pointer to the object structure containing information to access the object in future calls.

Signature:

```
void *(*open)(void *obj, H5VL_loc_params_t loc_params,
              H5I_type_t *opened_type, hid_t dxpl_id, void **req);
```

Arguments:

<code>obj</code>	(IN): Pointer to a file or group where the object needs to be opened or where the look-up of the target object needs to start.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>opened_type</code>	(OUT): buffer to return the type of the object opened (<code>H5I_GROUP</code> or <code>H5I_DATASET</code> or <code>H5I_DATATYPE</code>).
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `copy` callback in the object class should copy the object from the source object to the destination object. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*copy)(void *src_obj, H5VL_loc_params_t loc_params1,
               const char *src_name, void *dst_obj,
               H5VL_loc_params_t loc_params2, const char *dst_name,
               hid_t ocpypl_id, hid_t lcpl_id, hid_t dxpl_id, void **req);
```

Arguments:

<code>src_obj</code>	(IN): Pointer to location of the source object to be copied.
<code>loc_params1</code>	(IN): The location parameters as explained in section 2.1. The type should only be <code>H5VL_OBJECT_BY_SELF</code> for this callback.
<code>src_name</code>	(IN): Name of the source object to be copied.
<code>dst_obj</code>	(IN): Pointer to location of the destination object.
<code>loc_params2</code>	(IN): The location parameters as explained in section 2.1. The type should only be <code>H5VL_OBJECT_BY_SELF</code> for this callback.
<code>dst_name</code>	(IN): Name to be assigned to the new copy.
<code>ocpypl_id</code>	(IN): The object copy property list.
<code>lcpl_id</code>	(IN): The link creation property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the object class should retrieve information about the object as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *obj, H5VL_loc_params_t loc_params,
```

```
H5VL_object_get_t get_type, hid_t dxpl_id,
void **req, va_list arguments);
```

The `get_type` argument is an `enum`:

```
r object GET callback */
typedef enum H5VL_object_get_t {
    H5VL_REF_GET_NAME,          /* object name          */
    H5VL_REF_GET_REGION,       /* dataspace of region  */
    H5VL_REF_GET_TYPE          /* type of object        */
} H5VL_object_get_t;
```

Arguments:

<code>obj</code>	(IN): A location object where information needs to be retrieved from.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>get_type</code>	(IN): The type of the information to retrieve.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `get_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_REF_GET_NAME`, to retrieve a name for a referenced object:
 1. `ssize_t *ret` (OUT): buffer to return the length of the name.
 2. `char* name` (OUT): buffer to copy the name into.
 3. `size_t size` (IN): size of the buffer name, if 0, return only the buffer size needed.
 4. `H5R_type_t ret_type` (IN): type of region reference to query.
 5. `void *ref` (IN): the region reference to query.
- `H5VL_REF_GET_REGION`, to retrieve a region reference contained in `obj`:
 1. `hid_t *ret_id` (OUT): buffer for the dataspace created from the region reference.
 2. `H5R_type_t ret_type` (IN): type of region reference (should be `H5R_DATASET_REGION`).
 3. `void *ref` (IN): the region reference to open.
- `H5VL_REF_GET_TYPE`, to retrieve object type a reference points to:
 1. `H5O_type_t *type` (OUT): buffer to return the object type.
 2. `H5R_type_t ret_type` (IN): type of region reference to query.
 3. `void *ref` (IN): the region reference to query.

The `specific` callback in the object class implements specific operations on HDF5 objects as specified in the `specific_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,
                    H5VL_object_specific_t specific_type, hid_t dxpl_id, void
                    **req, va_list arguments);
```

The `specific_type` argument is an enum:

```
/* types for object SPECIFIC callback */
typedef enum H5VL_object_specific_t {
    H5VL_OBJECT_CHANGE_REF_COUNT, /* H5Oincr/decr_refcount */
    H5VL_OBJECT_EXISTS,          /* H5Oexists_by_name */
    H5VL_OBJECT_VISIT,           /* H5Ovisit(_by_name) */
    H5VL_REF_CREATE,             /* H5Rcreate */
} H5VL_object_specific_t;
```

Arguments:

<code>obj</code>	(IN): The location object where the operation needs to happen.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1.
<code>specific_type</code>	(IN): The type of the operation.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `specific_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_OBJECT_CHANGE_REF_COUNT`, to update the reference count for the object in `obj`:
 1. `int ref_count` (IN): reference count to set on the object.
- `H5VL_OBJECT_EXISTS`, to check if an object with name specified in `loc_params` (type `H5VL_OBJECT_BY_NAME`) exists:
 1. `htri_t *ret` (OUT): existence result, 0 if false, 1 if true.
- `H5VL_OBJECT_VISIT`, to iterate over all objects starting at the location provided and call a user specified callback on each one:
 1. `H5_index_t idx_type` (IN): Type of index.
 2. `H5_iter_order_t order` (IN): Order in which to iterate over index.
 3. `H5O_iterate_t op` (IN): User-defined function to pass each object to.

4. `void *op_data` (IN/OUT): User data to pass through to and to be returned by iterator operator function.

- `H5VL_REF_CREATE`, to create a reference of an object under the location object `obj`:

1. `void *ref` (OUT): the region reference created.
2. `char* name` (IN): Name of the object at the location `obj`.
3. `H5R_type_t ret_type` (IN): type of region reference to create.
4. `hid_t* space_id` (IN): Dataspace identifier with selection. Used only for dataset region references; passed as -1 if reference is an object reference, i.e., of type `H5R_OBJECT`.

The `optional` callback in the object class implements plugin specific operations on an HDF5 object. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
arguments);
```

Arguments:

- | | |
|------------------------|--|
| <code>obj</code> | (IN): The container or object where the operation needs to happen. |
| <code>dxpl_id</code> | (IN): The data transfer property list. |
| <code>req</code> | (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin. |
| <code>arguments</code> | (IN/OUT): argument list containing parameters and output pointers for the get operation. |

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

2.8 The Link Function Callbacks

The link API routines (H5L) allow HDF5 users to create and manage HDF5 links. All the H5L API routines that modify the HDF5 container map to one of the link callback routines in this class that the plugin needs to implement:

```
typedef struct H5VL_link_class_t {
    herr_t (*create)(H5VL_link_create_type_t create_type, void *obj,
        H5VL_loc_params_t loc_params, hid_t lcpl_id, hid_t lapl_id,
        hid_t dxpl_id, void **req);
    herr_t (*copy)(void *src_obj, H5VL_loc_params_t loc_params1, void
        *dst_obj, H5VL_loc_params_t loc_params2, hid_t lcpl, hid_t lapl,
        hid_t dxpl_id, void **req);
    herr_t (*move)(void *src_obj, H5VL_loc_params_t loc_params1, void
        *dst_obj, H5VL_loc_params_t loc_params2, hid_t lcpl, hid_t lapl,
        hid_t dxpl_id, void **req);
```

```

herr_t (*get)(void *obj, H5VL_loc_params_t loc_params,
              H5VL_link_get_t get_type, hid_t dxpl_id, void **req, va_list
              arguments);
herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,
                  H5VL_link_specific_t specific_type, hid_t dxpl_id, void **req,
                  va_list arguments);
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
                  arguments);
} H5VL_link_class_t;

```

The `create` callback in the group class should create a hard, soft, external, or user-defined links in the container. It returns an `herr_t` indicating success or failure.

Signature:

```

herr_t (*create)(H5VL_link_create_type_t create_type, void *obj,
                H5VL_loc_params_t loc_params, hid_t lcpl_id,
                hid_t lapl_id, hid_t dxpl_id, void **req);

```

The `create_type` argument is an enum:

```

/* link create types for VOL */
typedef enum H5VL_link_create_type_t {
    H5VL_LINK_CREATE_HARD, /* Hard Link */
    H5VL_LINK_CREATE_SOFT, /* Soft Link */
    H5VL_LINK_CREATE_UD /* External / UD Link */
} H5VL_link_create_type_t;

```

Arguments:

<code>create_type</code>	(IN): type of the link to be created.
<code>obj</code>	(IN): Pointer to an object where the link needs to be created from.
<code>loc_params</code>	(IN): The location parameters as explained in section 2.1 for the source object.
<code>lcpl_id</code>	(IN): The link creation property list. It contains all the link creation properties in addition to other API parameters depending on the creation type, which will be detailed next.
<code>lapl_id</code>	(IN): The link access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

As mentioned in the argument list, the `lcpl_id` contains the parameters for the link creation operation depending on the creation type:

- `H5VL_LINK_CREATE_HARD` contains two properties:
 1. `H5VL_LINK_TARGET` (with type `void*`): The target object where the hard link needs to be created to.
 2. `H5VL_LINK_TARGET_LOC_PARAMS` (with type `H5VL_loc_params_t`): The location parameters as explained in section 2.1 for the target object.

- `H5VL_LINK_CREATE_SOFT` contains one property:
 1. `H5VL_LINK_TARGET_NAME` (with type `char*`): The target link where the soft link should point to.
- `H5VL_LINK_CREATE_UD` contains two properties:
 1. `H5VL_LINK_TYPE` (with type `H5L_type_t`): The user defined link class. `H5L_TYPE_EXTERNAL` suggests an external link is to be created.
 2. `H5VL_LINK_UDATA` (with type `void*`): User supplied link information (contains the external link buffer for external links).
 3. `H5VL_LINK_UDATA_SIZE` (with type `size_t`): size of the `udata` buffer.

The move callback in the link class should move a link within the HDF5 container. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*move)(void *src_obj, H5VL_loc_params_t loc_params1, void
               *dst_obj, H5VL_loc_params_t loc_params2, hid_t lcpl, hid_t lapl,
               hid_t dxpl_id, void **req);
```

Arguments:

<code>src_obj</code>	(IN): original/source object or file.
<code>loc_params1</code>	(IN): The location parameters for the source object as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_NAME</code> in this callback.
<code>dst_obj</code>	(IN): destination object or file.
<code>loc_params1</code>	(IN): The location parameters for the destination object as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_NAME</code> in this callback.
<code>lcpl_id</code>	(IN): The link creation property list.
<code>lapl_id</code>	(IN): The link access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The copy callback in the link class should copy a link within the HDF5 container. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*copy)(void *src_obj, H5VL_loc_params_t loc_params1, void
               *dst_obj, H5VL_loc_params_t loc_params2, hid_t lcpl, hid_t lapl,
               hid_t dxpl_id, void **req);
```

Arguments:

<code>src_obj</code>	(IN): original/source object or file.
<code>loc_params1</code>	(IN): The location parameters for the source object as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_NAME</code> in this callback.
<code>dst_obj</code>	(IN): destination object or file.
<code>loc_params1</code>	(IN): The location parameters for the destination object as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_NAME</code> in this callback.
<code>lcpl_id</code>	(IN): The link creation property list.
<code>lapl_id</code>	(IN): The link access property list.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.

The `get` callback in the link class should retrieve information about links as specified in the `get_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*get)(void *obj, H5VL_loc_params_t loc_params,
              H5VL_link_get_t get_type, hid_t dxpl_id, void **req,
              va_list arguments);
```

The `get_type` argument is an enum:

```
/* types for all link get API routines */
typedef enum H5VL_link_get_t {
    H5VL_LINK_GET_INFO,      /* link info      */
    H5VL_LINK_GET_NAME,      /* link name      */
    H5VL_LINK_GET_VAL        /* link value     */
} H5VL_link_get_t;
```

Arguments:

<code>obj</code>	(IN): The file or group object where information needs to be retrieved from.
<code>loc_params</code>	(IN): The location parameters for the source object as explained in section 2.1. The type can be only <code>H5VL_OBJECT_BY_NAME</code> or <code>H5VL_OBJECT_BY_IDX</code> in this callback.
<code>get_type</code>	(IN): The type of the information to retrieve.
<code>dxpl_id</code>	(IN): The data transfer property list.
<code>req</code>	(IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
<code>arguments</code>	(IN/OUT): argument list containing parameters and output pointers for the get operation.

The `arguments` argument contains a variable list of arguments depending on the `get_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_LINK_GET_INFO`, to retrieve the link info from the link specified in the `loc_params`:
 1. `H5L_info_t info` (OUT): pointer to info structure to fill.
- `H5VL_LINK_GET_NAME`, to retrieve the name of the link specified by the index information in `loc_params` (`loc_params` is of type `H5VL_OBJECT_BY_IDX` only with this type):
 1. `char* name` (OUT): buffer to copy the name into.
 2. `size_t size` (IN): size of the buffer name, if 0, return only the buffer size needed.
 3. `ssize_t *ret` (OUT): buffer to return the length of the link name.
- `H5VL_LINK_GET_VAL`, to retrieve the link value from the link specified in the `loc_params`:
 1. `void *buf` (OUT): buffer to put the value into.
 2. `size_t size` (IN): size of the passed in buffer.

The `specific` callback in the link class implements specific operations on HDF5 links as specified in the `specific_type` parameter. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*specific)(void *obj, H5VL_loc_params_t loc_params,
                  H5VL_link_specific_t specific_type, hid_t dxpl_id, void
                  **req, va_list arguments);
```

The `specific_type` argument is an enum:

```
/* types for link SPECIFIC callback */
typedef enum H5VL_link_specific_t {
    H5VL_LINK_DELETE,      /* H5Ldelete(_by_idx)          */
    H5VL_LINK_EXISTS,      /* link existence              */
    H5VL_LINK_ITER         /* H5Literate/visit(_by_name) */
} H5VL_link_specific_t;
```

Arguments:

- | | |
|----------------------------|--|
| <code>obj</code> | (IN): The location object where the operation needs to happen. |
| <code>loc_params</code> | (IN): The location parameters as explained in section 2.1. |
| <code>specific_type</code> | (IN): The type of the operation. |
| <code>dxpl_id</code> | (IN): The data transfer property list. |
| <code>req</code> | (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin. |
| <code>arguments</code> | (IN/OUT): argument list containing parameters and output pointers for the get operation. |

The `arguments` argument contains a variable list of arguments depending on the `specific_type` parameter. The following list shows the argument list, in-order, for each type:

- `H5VL_LINK_DELETE`, to remove a link specified in the location parameter from an HDF5 container.
- `H5VL_LINK_EXISTS`, to determine whether the link specified in the `loc_params` exists (`loc_params` is of type `H5VL_OBJECT_BY_NAME` only with this type):
 1. `htri_t *ret` (OUT): buffer for the existence of the link (0 for no, 1 for yes).
- `H5VL_LINK_ITER`, to iterate over all links starting at the location provided and call a user specified callback on each one:
 1. `hbool_t recursive` (IN): whether to recursively follow links into subgroups of the specified group.
 2. `H5_index_t idx_type` (IN): Type of index.
 3. `H5_iter_order_t order` (IN): Order in which to iterate over index.
 4. `hsize_t *idx` (IN/OUT): iteration position where to start and return position where an interrupted iteration may restart.
 5. `H5L_iterate_t op` (IN): User-defined function to pass each link to.
 6. `void *op_data` (IN/OUT): User data to pass through to and to be returned by iterator operator function.

The `optional` callback in the link class implements plugin specific operations on an HDF5 link. It returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
arguments);
```

Arguments:

- | | |
|------------------------|--|
| <code>obj</code> | (IN): The container or object where the operation needs to happen. |
| <code>dxpl_id</code> | (IN): The data transfer property list. |
| <code>req</code> | (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin. |
| <code>arguments</code> | (IN/OUT): argument list containing parameters and output pointers for the get operation. |

Each plugin should be able to parse the `va_list arguments` if it has plugin specific operations to implement and determine the type of the operation and the parameters through a predefined schema.

2.9 The Asynchronous Function Callbacks

As of now, the HDF5 library does not provide asynchronous API operations. An asynchronous class to manage asynchronous operations was added nevertheless to handle an asynchronous API that might be added in the future:

```
typedef struct H5VL_async_class_t {  
    herr_t (*cancel)(void **, H5ES_status_t *);  
    herr_t (*test) (void **, H5ES_status_t *);  
    herr_t (*wait) (void **, H5ES_status_t *);  
} H5VL_async_class_t;
```

The `H5ES_status_t` argument is an enum:

```
/* Asynchronous operation status */  
typedef enum H5ES_status_t {  
    H5ES_STATUS_IN_PROGRESS, /* Operation has not yet completed */  
    H5ES_STATUS_SUCCEED,    /* Operation has completed, successfully */  
    H5ES_STATUS_FAIL,      /* Operation has completed, but failed */  
    H5ES_STATUS_CANCEL     /* Operation has not completed and has been  
                           cancelled */  
} H5ES_status_t;
```

The `cancel` callback attempts to cancel an asynchronous operation and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*cancel)(void **req, H5ES_status_t *status);
```

Arguments:

`req` (IN): A pointer to the asynchronous request.
`status` (OUT): result of the cancel operation.

The `test` callback tests an asynchronous operation completion and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*test)(void **req, H5ES_status_t *status);
```

Arguments:

`req` (IN): A pointer to the asynchronous request.
`status` (OUT): result of the test operation.

The `wait` callback waits for an asynchronous operation completion and returns an `herr_t` indicating success or failure.

Signature:

```
herr_t (*wait)(void **req, H5ES_status_t *status);
```

Arguments:

req (IN): A pointer to the asynchronous request.
status (OUT): result of the wait operation.

2.10 The Optional Generic Callback

A generic optional callback is provided for services that are specific to a plugin. Asynchronous operations/services could have been implemented through this optional callback, however since we anticipate asynchronous operations to be common across many plugins, they have been promoted to a subclass by themselves as shown in the previous section. Other services that are not shared across different plugins might include transactional semantics for managing operations like in /reffastforward, MSC - add more use cases here.

The **optional** callback has the following definition. It returns an **herr_t** indicating success or failure.

Signature:

```
herr_t (*optional)(void *obj, hid_t dxpl_id, void **req, va_list
    arguments);
```

Arguments:

obj (IN): The container or object where the operation needs to happen.
dxpl_id (IN): The data transfer property list.
req (IN/OUT): A pointer to the asynchronous request of the operation created by the plugin.
arguments (IN/OUT): argument list containing parameters and output pointers for the get operation.

3 New VOL API Routines

New API routines have been added to the HDF5 library to manage VOL plugins. This section details each new API call and explains their intended usage.

```
hid_t H5VLregister(const H5VL_class_t *cls);
```

Registers a user defined VOL plugin with the HDF5 library and returns an identifier for that plugin to be used if successful, otherwise returns a negative value. This function is used when the application has direct access to the plugin it wants to use and is able to obtain a pointer for the plugin structure to pass to the HDF5 library. The plugin must be unregistered before the application exists with **H5VLunregister()**.

Arguments:

cls (IN): A pointer to the plugin structure to register.

```
hid_t H5VLregister_by_name(const char *plugin_name);
```

Registers a VOL plugin with the HDF5 library given the name of the plugin and returns an identifier for that plugin to be used if successful, otherwise returns a negative value. If the plugin is already registered the library will create an identifier for it and returns it to the user; otherwise the library will look for available third party plugins and dynamically load them and registers them for use, then returns an identifier for the plugin. Third party plugins must be installed on the system as a shared library or DLL. The plugin must be unregistered before the application exists with `H5VLunregister()`.

Arguments:

`name` (IN): The plugin name to search for an register.

```
herr_t H5VLunregister(hid_t plugin_id);
```

Unregisters a plugin from the library and return a positive value on success otherwise return a negative value.

Arguments:

`plugin_id` (IN): A valid identifier of the plugin to unregister.

```
htri_t H5VLis_registered(const char *name);
```

Checks if a VOL plugin is registered with the library given the plugin name and returns TRUE/FALSE on success, otherwise it returns a negative value.

Arguments:

`name` (IN): The plugin name to check for.

```
ssize_t H5VLget_plugin_name(hid_t obj_id, char *name/*out*/, size_t size);
```

Retrieves the name of a VOL plugin given and object identifier that was created/opened with it. On success, the name length is returned.

Arguments:

`obj_id` (IN): The object identifier to check.

`name` (OUT): Buffer pointer to put the plugin name. If NULL, the library just returns the size required to store the plugin name.

`size` (IN): the size of the passed in buffer.

```
hid_t H5VLget_plugin_id(const char *name);
```

Given a plugin name that is registered with the library, this function returns an identifier for the plugin. If the plugin is not registered with the library, a negative value is returned. The identifier must be released with a call to `H5VLclose()`.

Arguments:

`name` (IN): The plugin name to check for.

```
herr_t H5VLclose(hid_t plugin_id);
```

Shuts down access to the plugin that the identifier points to and release resources associated with it.

Arguments:

`plugin_id` (IN): A valid identifier of the plugin to close.

```
hid_t H5VLObject_register(void *obj, H5I_type_t obj_type, hid_t
    plugin_id);
```

Creates an HDF5 identifier for an object given its type and an identifier of the plugin that was used in creating it. This function is typically used by external plugins wanting to wrap an `hid_t` around a plugin specific object to return to a user callback for functions like `H5Literate`, `H5Lvisit`, `H5Ovisit`, `H5Aiterate` etc... A Negative value is returned on failure.

Arguments:

`obj` (IN): pointer to the object to register.

`obj_type` (IN): The type of the object (`H5I_FILE`, `H5I_DATASET`, `H5I_GROUP`, `H5I_ATTRIBUTE`, `H5I_DATATYPE`).

`plugin_id` (IN): identifier of the vol plugin the object is associated with.

```
herr_t H5VLget_object(hid_t obj_id, void **obj);
```

Retrieves a pointer to the VOL object from an HDF5 file or object identifier.

Arguments:

`obj_id` (IN): identifier of the object to dereference.

`obj` (OUT): buffer to store the pointer of the VOL object associated with the object identifier.

```
herr_t H5VLinitialize(hid_t plugin_id, hid_t vipl_id);
```

Initializes access to the VOL plugin represented by `plugin_id` with initialization parameters set as properties in the the property list `vipl_id`. Some plugins may not require initialization. For more information about plugin initialization and what properties are required to be set in the property list refer to the actual plugin documentation.

Arguments:

`plugin_id` (IN): identifier of VOL plugin to initialize.

`vipl_id` (IN): VOL initialize property list containing properties required for plugin initialization.

```
herr_t H5VLterminate(hid_t plugin_id, hid_t vtpl_id);
```

Terminates access to the VOL plugin represented by `plugin_id` with termination parameters set as properties in the the property list `vtpl_id`. Some plugins may not require termination. For more information about plugin termination and what properties are required to be set in the property list refer to the actual plugin documentation.

Arguments:

- plugin_id (IN): identifier of VOL plugin to terminate.
- vtpl_id (IN): VOL terminate property list containing properties required for plugin termination.

A set of API calls that map directly to the VOL callbacks themselves have been added to aid in the development of stacked or mirrored plugins. For more information on usage refer to Section 6. The routines are listed below.

```

/* ATTRIBUTE OBJECT ROUTINES */
void *H5VLattr_create(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *attr_name, hid_t acpl_id, hid_t aapl_id,
    hid_t dxpl_id, void **req);
void *H5VLattr_open(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t aapl_id, hid_t dxpl_id, void
    **req);
herr_t H5VLattr_read(void *attr, hid_t plugin_id, hid_t dtype_id, void
    *buf, hid_t dxpl_id, void **req);
herr_t H5VLattr_write(void *attr, hid_t plugin_id, hid_t dtype_id, const
    void *buf, hid_t dxpl_id, void **req);
herr_t H5VLattr_get(void *obj, hid_t plugin_id, H5VL_attr_get_t
    get_type, hid_t dxpl_id, void **req, va_list arguments);
herr_t H5VLattr_specific(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, H5VL_attr_specific_t specific_type, hid_t dxpl_id, void
    **req, va_list arguments);
herr_t H5VLattr_optional(void *obj, hid_t plugin_id, hid_t dxpl_id, void
    **req, va_list arguments);
herr_t H5VLattr_close(void *attr, hid_t plugin_id, hid_t dxpl_id, void
    **req);

/* DATASET OBJECT ROUTINES */
void *H5VLdataset_create(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t dcpl_id, hid_t dapl_id, hid_t
    dxpl_id, void **req);
void *H5VLdataset_open(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t dapl_id, hid_t dxpl_id, void
    **req);
herr_t H5VLdataset_read(void *dset, hid_t plugin_id, hid_t mem_type_id,
    hid_t mem_space_id, hid_t file_space_id, hid_t plist_id, void *buf,
    void **req);
herr_t H5VLdataset_write(void *dset, hid_t plugin_id, hid_t mem_type_id,
    hid_t mem_space_id, hid_t file_space_id, hid_t plist_id, const void
    *buf, void **req);
herr_t H5VLdataset_get(void *dset, hid_t plugin_id, H5VL_dataset_get_t
    get_type, hid_t dxpl_id, void **req, va_list arguments);
herr_t H5VLdataset_specific(void *obj, hid_t plugin_id,
    H5VL_dataset_specific_t specific_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLdataset_optional(void *obj, hid_t plugin_id, hid_t dxpl_id,
    void **req, va_list arguments);
herr_t H5VLdataset_close(void *dset, hid_t plugin_id, hid_t dxpl_id,
    void **req);

```

```

/* DATATYPE OBJECT ROUTINES */
void *H5VLdatatype_commit(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t type_id, hid_t lcpl_id, hid_t
    tcpl_id, hid_t tapl_id, hid_t dxpl_id, void **req);
void *H5VLdatatype_open(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t tapl_id, hid_t dxpl_id, void
    **req);
herr_t H5VLdatatype_get(void *dt, hid_t plugin_id, H5VL_datatype_get_t
    get_type, hid_t dxpl_id, void **req, va_list arguments);
herr_t H5VLdatatype_specific(void *obj, hid_t plugin_id,
    H5VL_datatype_specific_t specific_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLdatatype_optional(void *obj, hid_t plugin_id, hid_t dxpl_id,
    void **req, va_list arguments);
herr_t H5VLdatatype_close(void *dt, hid_t plugin_id, hid_t dxpl_id, void
    **req);

/* FILE OBJECT ROUTINES */
void *H5VLfile_create(const char *name, unsigned flags, hid_t fcpl_id,
    hid_t fapl_id, hid_t dxpl_id, void **req);
void *H5VLfile_open(const char *name, unsigned flags, hid_t fapl_id,
    hid_t dxpl_id, void **req);
herr_t H5VLfile_get(void *file, hid_t plugin_id, H5VL_file_get_t
    get_type, hid_t dxpl_id, void **req, va_list arguments);
herr_t H5VLfile_specific(void *obj, hid_t plugin_id,
    H5VL_file_specific_t specific_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLfile_optional(void *obj, hid_t plugin_id, hid_t dxpl_id, void
    **req, va_list arguments);
herr_t H5VLfile_close(void *file, hid_t plugin_id, hid_t dxpl_id, void
    **req);

/* GROUP OBJECT ROUTINES */
void *H5VLgroup_create(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t gcpl_id, hid_t gapl_id, hid_t
    dxpl_id, void **req);
void *H5VLgroup_open(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, const char *name, hid_t gapl_id, hid_t dxpl_id, void
    **req);
herr_t H5VLgroup_get(void *obj, hid_t plugin_id, H5VL_group_get_t
    get_type, hid_t dxpl_id, void **req, va_list arguments);
herr_t H5VLgroup_specific(void *obj, hid_t plugin_id,
    H5VL_group_specific_t specific_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLgroup_optional(void *obj, hid_t plugin_id, hid_t dxpl_id,
    void **req, va_list arguments);
herr_t H5VLgroup_close(void *grp, hid_t plugin_id, hid_t dxpl_id, void
    **req);

/* LINK OBJECT ROUTINES */

```

```

herr_t H5VLlink_create(H5VL_link_create_type_t create_type, void *obj,
    H5VL_loc_params_t loc_params, hid_t plugin_id, hid_t lcpl_id, hid_t
    lapl_id, hid_t dxpl_id, void **req);
herr_t H5VLlink_copy(void *src_obj, H5VL_loc_params_t loc_params1, void
    *dst_obj, H5VL_loc_params_t loc_params2, hid_t plugin_id, hid_t
    lcpl_id, hid_t lapl_id, hid_t dxpl_id, void **req);
herr_t H5VLlink_move(void *src_obj, H5VL_loc_params_t loc_params1, void
    *dst_obj, H5VL_loc_params_t loc_params2, hid_t plugin_id, hid_t
    lcpl_id, hid_t lapl_id, hid_t dxpl_id, void **req);
herr_t H5VLlink_get(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, H5VL_link_get_t get_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLlink_specific(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, H5VL_link_specific_t specific_type, hid_t dxpl_id, void
    **req, va_list arguments);
herr_t H5VLlink_optional(void *obj, hid_t plugin_id, hid_t dxpl_id, void
    **req, va_list arguments);

/* OBJECT ROUTINES */
void *H5VLobject_open(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, H5L_type_t *opened_type, hid_t dxpl_id, void **req);
herr_t H5VLobject_copy(void *src_obj, H5VL_loc_params_t loc_params1,
    hid_t plugin_id1, const char *src_name, void *dst_obj,
    H5VL_loc_params_t loc_params2, hid_t plugin_id2, const char
    *dst_name, hid_t ocpypl_id, hid_t lcpl_id, hid_t dxpl_id, void
    **req);
herr_t H5VLobject_get(void *obj, H5VL_loc_params_t loc_params, hid_t
    plugin_id, H5VL_object_get_t get_type, hid_t dxpl_id, void **req,
    va_list arguments);
herr_t H5VLobject_specific(void *obj, H5VL_loc_params_t loc_params,
    hid_t plugin_id, H5VL_object_specific_t specific_type, hid_t
    dxpl_id, void **req, va_list arguments);
herr_t H5VLobject_optional(void *obj, hid_t plugin_id, hid_t dxpl_id,
    void **req, va_list arguments);

/* ASYNCHRONOUS ROUTINES */
herr_t H5VLrequest_cancel(void **req, hid_t plugin_id, H5ES_status_t
    *status);
herr_t H5VLrequest_test(void **req, hid_t plugin_id, H5ES_status_t
    *status);
herr_t H5VLrequest_wait(void **req, hid_t plugin_id, H5ES_status_t
    *status);

```

4 Creating and Using an Internal Plugin

Internal plugins are developed internally with the HDF5 library and are required to ship with the entire library to be used. Typically those plugins need to use internal features and functions of the HDF5 library that are not available publicly from the user application.

4.1 Implementing an Internal Plugin

The first step to implement an internal plugin is to implement all the callbacks defined in the VOL class as described in section 2. After implementing the VOL class, the next step would be to allow users to select this plugin to be used. This is done by creating a new API routine to set the plugin on the file access property list. For example, if we create an internal plugin called “dummy” that needs an MPI communicator and info object as information from the user, that routine signature should look like:

```
herr_t H5Pset_fapl_dummy(hid_t fapl_id, MPI_Comm comm, MPI_Info info);
```

The implementation for the above routine should use the internal function:

```
herr_t H5P_set_vol(H5P_genplist_t *plist, H5VL_class_t *vol_cls, const
void *vol_info);
```

that will set the file access using that `fapl_id` to go through the “dummy” plugin. It will also call the copy callback of the “dummy” plugin on the info object (`comm` and `info`).

A sample implementation for the `H5Pset_fapl_dummy()` would look like this:

```
/* DUMMY-specific file access properties */
typedef struct H5VL_dummy_fapl_t {
    MPI_Comm      comm; /* communicator */
    MPI_Info      info; /* MPI information */
} H5VL_dummy_fapl_t;

herr_t
H5Pset_fapl_dummy(hid_t fapl_id, MPI_Comm comm, MPI_Info info)
{
    H5VL_dummy_fapl_t fa;
    H5P_genplist_t *plist; /* Property list pointer */
    herr_t          ret_value;

    FUNC_ENTER_API(FAIL)

    if(fapl_id == H5P_DEFAULT)
        HGOTO_ERROR(H5E_PLIST, H5E_BADVALUE, FAIL, "can't set values in
        default property list")

    if(NULL == (plist = H5P_object_verify(fapl_id, H5P_FILE_ACCESS)))
        HGOTO_ERROR(H5E_ARGS, H5E_BADTYPE, FAIL, "not a file access
        property list")

    if(MPI_COMM_NULL == comm)
        HGOTO_ERROR(H5E_PLIST, H5E_BADTYPE, FAIL, "not a valid
        communicator")

    /* Initialize driver specific properties */
    fa.comm = comm;
```

```

fa.info = info;

ret_value = H5P_set_vol(plist, &H5VL_dummy_g, &fa);

done:
    FUNC_LEAVE_API(ret_value)
} /* end H5Pset_fapl_dummy() */

```

At this point, the internal plugin is ready to be used. For more information on how to implement an internal plugin, the native plugin for the HDF5 library is a comprehensive plugin that implements all features of the library and can be used as a guide.

4.2 Using an Internal Plugin

An internal plugin is registered automatically with the HDF5 library when the library gets initialized. Users should not attempt to register an internal plugin. The internal registration operation creates a global identifier for each internal plugin. Applications can query the library for that identifier using the plugin name with `H5VLget_plugin_id()`. This identifier is used to initialize the plugin if it requires an initialization phase. A third party library could use that ID to create a plugin that stacks or mirrors on top of an internal plugin using that identifier (refer to section 6 for more details on stacking and mirroring).

Some plugins could require an initialization phase before using the plugin with some input parameters from users. For example, a remote plugin would need to setup a connection to a remote machine with a specific hostname that the user knows and the HDF5 library does not know about. This process is done by creating a VOL initialization property list and setting a public property that should be defined by the internal plugin to the hostname of the remote machine. The user then calls the initialize routine (`H5VLintialize()`) with the plugin_id and the property list to initialize the plugin. Some plugins could create a public API routine that wraps the property list creation, property settings, and plugin initialization into one routine that the user calls to simplify this process for the users. The native HDF5 plugin does not require an initialization phase.

Now that the plugin is initialized, the application can create or open an HDF5 container using the internal plugin by creating a file access property list and calling the corresponding API routine that the internal plugin should define for setting the VOL plugin property and container parameters on the file access property list. That FAPL can be passed to the `H5Fcreate()` or `H5Fopen` calls to access the container through the VOL plugin. Subsequently, all other HDF5 calls on the container and objects created or opened in that container will automatically go to the selected VOL plugin.

After closing all objects and containers on that plugin, the application should terminate access to the selected plugin if it requires a termination phase by calling `H5VLterminate()` with the required termination properties (similar to the initialization phase). The native HDF5 plugin does not require a termination phase.

The internal plugin is unregistered automatically by the HDF5 library on library termination, and users should not attempt to unregister it themselves.

5 Creating and Using an External Plugin

External plugins are developed outside of the HDF5 library and do not use any internal HDF5 private functions. They do not require to be shipped with the HDF5 library, but can just link to it from userspace just like an HDF5 application.

5.1 New API Routines for External Plugins

Some callbacks in the VOL class require new API routines for the implementation to be possible. Two new API routines have been added for that matter:

```
hid_t H5VLObject_register(void *obj, H5I_type_t obj_type, const
                          H5VL_class_t *cls);
```

to register an `hid_t` with an object `obj` associated with the VOL plugin of class `cls`. This is needed in all iterate and visit callbacks where the plugins internally need to wrap an `hid_t` around an object to call the user defined operation `op` on.

```
herr_t H5VLget_object(hid_t obj_id, void **obj, H5VL_t **vol_plugin);
```

to retrieve the VOL object and plugin structure from an HDF5 identifier (`hid_t`). The plugin structure is defined as:

```
struct H5VL_t {
    const H5VL_class_t *cls; /* constant class info */
    const char *container_name; /* name of the underlying storage
                                container */
    unsigned long feature_flags; /* VOL Driver feature Flags */
    int nrefs; /* number of references by objects
              using this struct */
};
```

5.2 Using an External Plugin

Unlike internal plugins, the external plugins cannot create an API routine for applications to use to set the VOL plugin in the file access property list. After implementing the VOL class as described in section 2, the application has to register the plugin with HDF5 library. The function to do that is `H5VLregister()`:

```
hid_t H5VLregister(const H5VL_class_t *cls);
```

where `cls` is a pointer to the external plugin to be used. The identifier returned can be used to set this plugin to be used in the file access property list with this API routine:

```
herr_t H5Pset_vol(hid_t fapl_id, hid_t plugin_id, const void  
                *new_vol_info);
```

where `plugin_id` is the identifier returned from the `H5VLregister()` and `new_vol_info` is the plugin information needed from the application.

The user is required to un-register the plugin from the library when access to the container(s) is terminated using:

```
herr_t H5VLunregister(hid_t plugin_id);
```

6 Interchanging and Stacking VOL Plugins

Accessing an HDF5 container with a VOL plugin different than the one it was created with would be a valid approach as long as the underlying file format is the same. This would be the user's responsibility to ensure that the different plugins are interchangeable.

6.1 Stacking Plugins on Top of Each Other

It would be also possible to stack VOL plugins on top of each other. This notion is similar to the idea of the split VFD, where underneath the split VFD itself, two file drivers would be used, one for the file storing the metadata and another for raw data. Some stackings make sense and others would be erroneous. For example, stacking the native HDF5 plugin on top of a non-HDF5 backend plugin does not make sense and is erroneous. Figure 1 shows a stacking of a remote plugin, where data is distributed remotely, on top of the native h5 plugin, where servers that store the data at remote locations use the h5 file format.

6.2 Mirroring Plugins

Another useful design option is to allow a mirroring plugin, where the HDF5 API calls are forwarded through a mirror plugin to two or more VOL plugins. This is an extension to the stacking feature. Figure 2 shows an example of a VOL mirror that maps HDF5 API calls to an h5 backend plugin and an XML backend plugin.

Another possible VOL plugin could be a statistics plugin that just gathers information on HDF5 API calls and records statistics associated with the number




Figure 1: Stacked VOL plugins.

of calls to a specific API functions and corresponding parameters. This plugin would be very useful for profiling purposes. The statistics plugin would be stacked on top of another VOL plugin that actually performs the required access to the file.

6.3 Implementing Stacked and Mirrored Plugins

A new set of API calls that map directly to the VOL callbacks have been added to the HDF5 library to make stacking and mirroring easy for plugin developers. Those APIs are listed at the end of Section 3. They are similar to the public

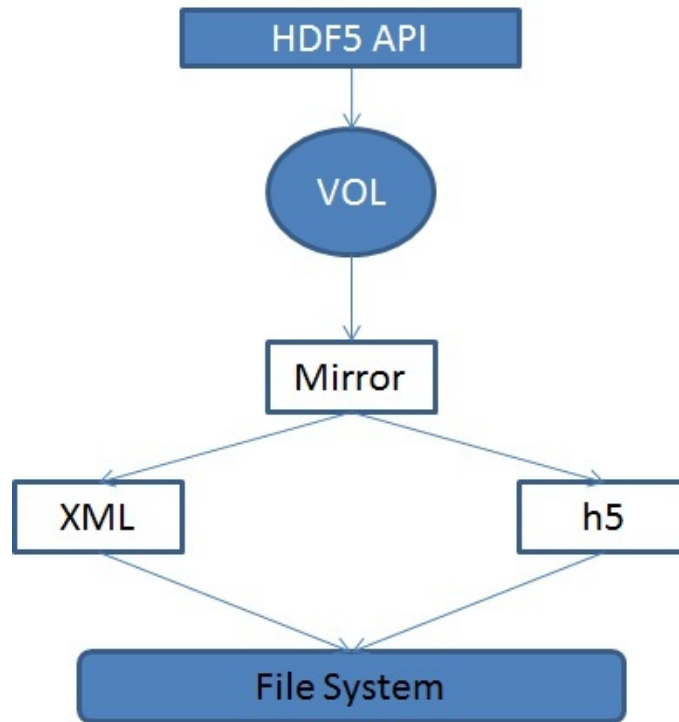


Figure 2: Mirrored VOL plugins.

VFD (H5FD) routines that call the VFD callbacks directly.

Using those new APIs, A mirror VOL plugin could, for example, implement the create callback in the file class by creating two HDF5 containers each with a different underlying plugin with 2 calls to:

```
void *H5VLfile_create(const char *name, unsigned flags, hid_t fcpl_id,
                    hid_t fapl_id, hid_t dxpl_id, void **req);
```

where each call uses a different file access property list, each with the VOL plugin property set to the underlying plugins. Then the mirror plugin creates a structure containing the pointers for the containers returned from both underlying plugins and returns a pointer to that structure as the output of the file create callback. Subsequent access to the container would again call the the public VOL routines twice for each plugin with its corresponding objects and combine the output as needed.